

EE382N (20): Computer Architecture - Parallelism and Locality
Lecture 6 – GPUs (I)

Mattan Erez



The University of Texas at Austin

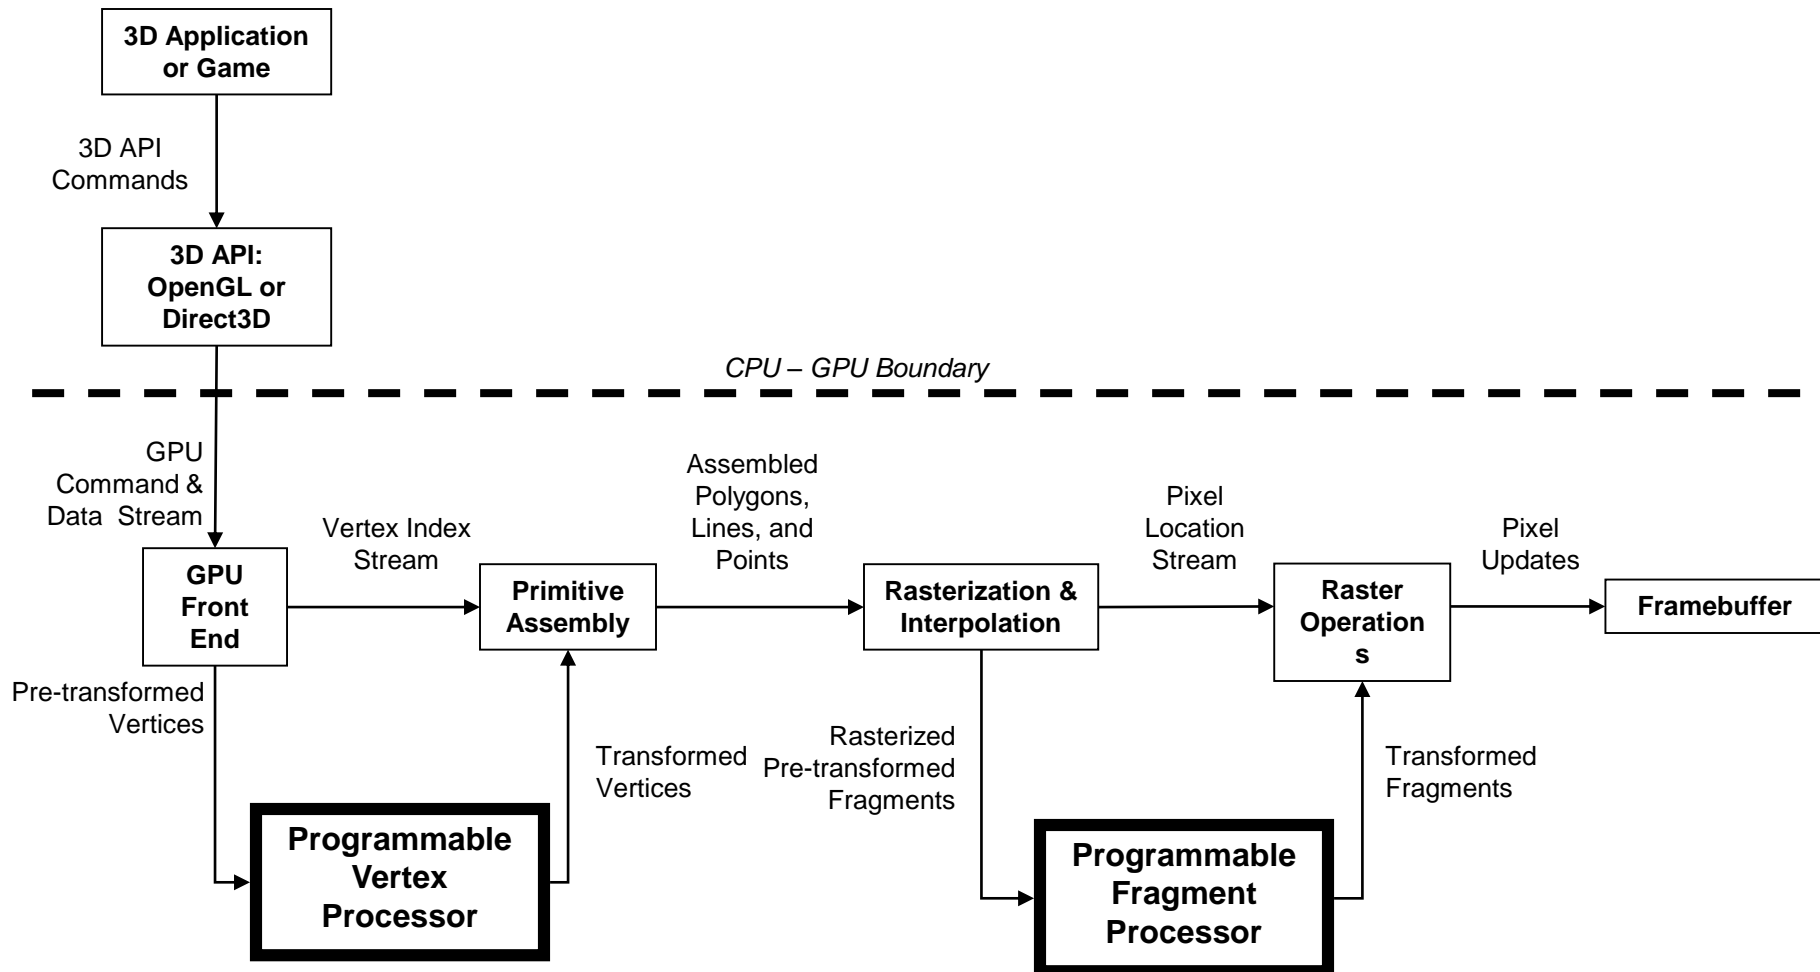


A GPU Renders 3D Scenes

- A **Graphics Processing Unit (GPU)** accelerates rendering of 3D scenes
 - Input: description of scene
 - Output: colored pixels to be displayed on a screen
- Input:
 - Geometry (triangles), colors, lights, effects, textures
- Output:



Adding Programmability to the Graphics Pipeline



Vertex and Fragment Processing Share Unified Processing Elements


4

- Load balancing HW is a problem


Vertex Shader



Pixel Shader




Vertex Shader



Pixel Shader



Heavy Geometry
Workload Perf = 4



Heavy Pixel
Workload Perf = 8



Vertex and Fragment Processing Share Unified Processing Elements


5

- Load balancing SW is easier


Unified Shader



Unified Shader



Heavy Geometry
Workload Perf = 11



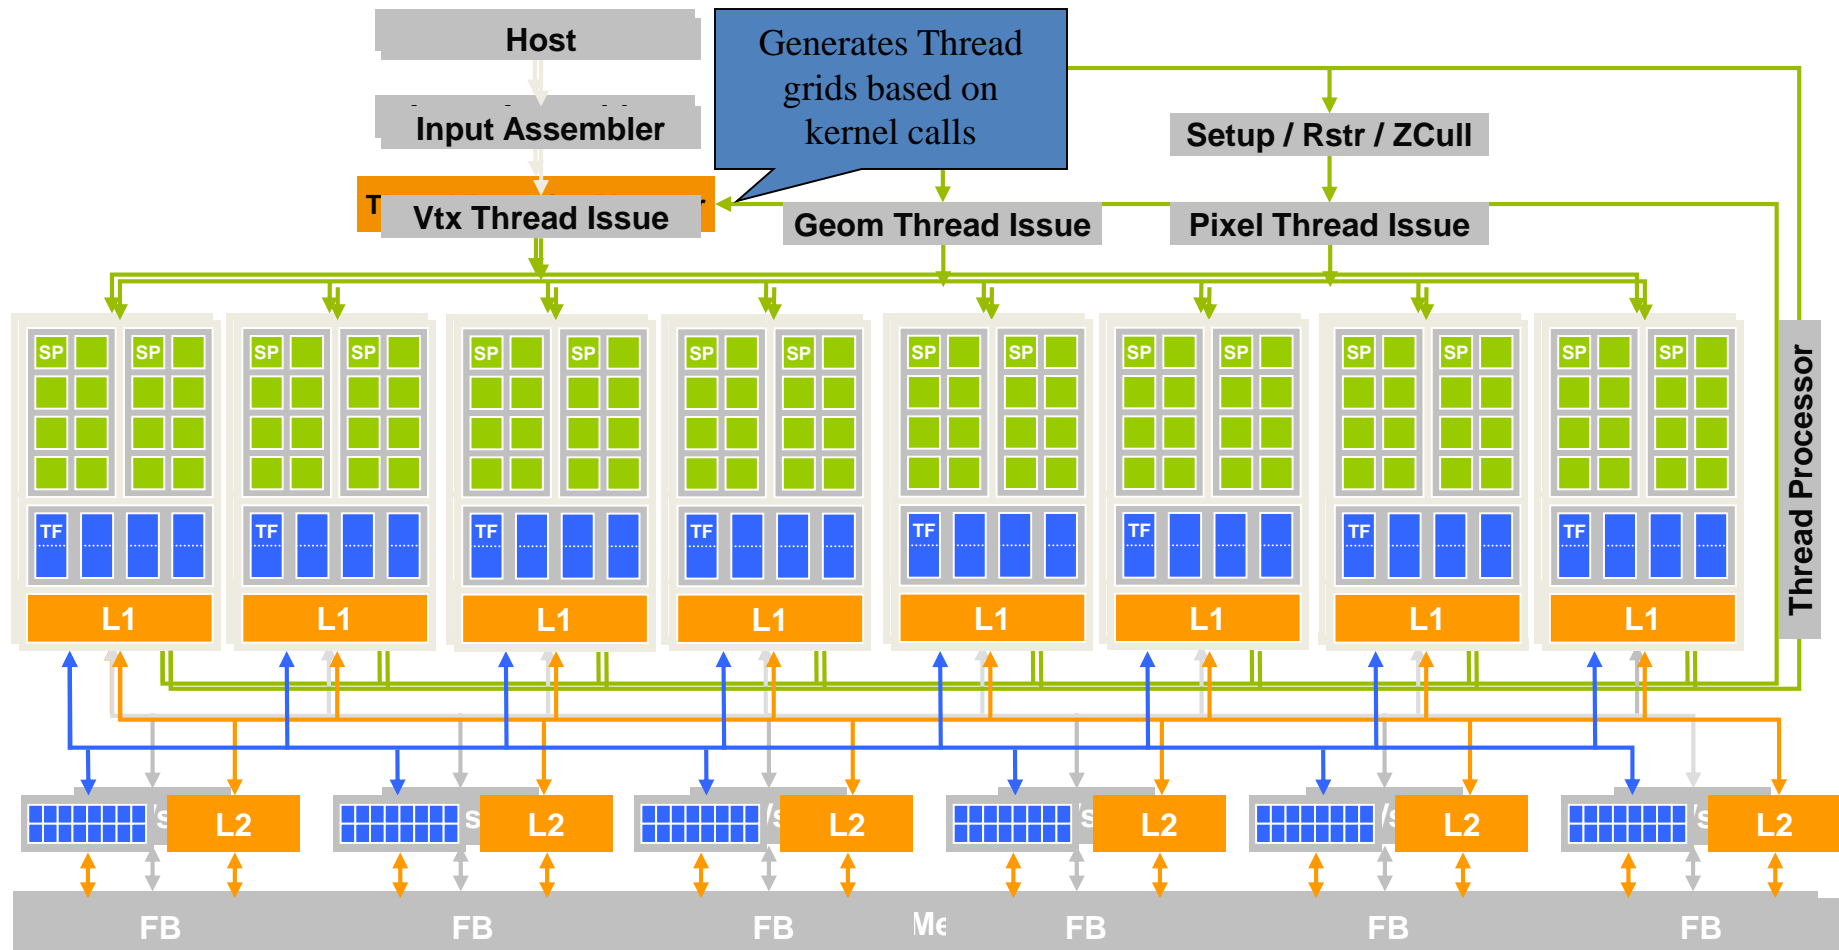
Heavy Pixel
Workload Perf = 11



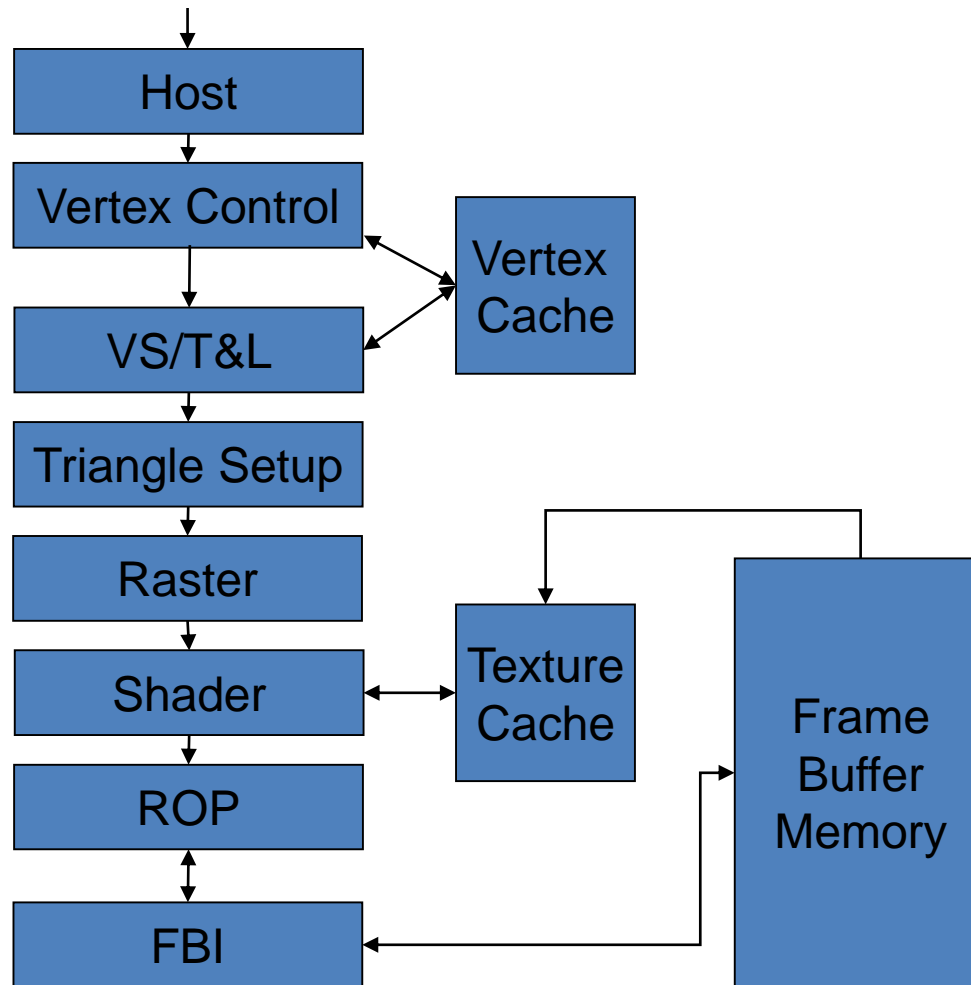
Make the Compute Core The Focus of the Architecture

6

- The future of GPUs is programmable processing
- Processors execute computing threads
- So – build the architecture around the processor
- Alternative operating mode specifically for computing

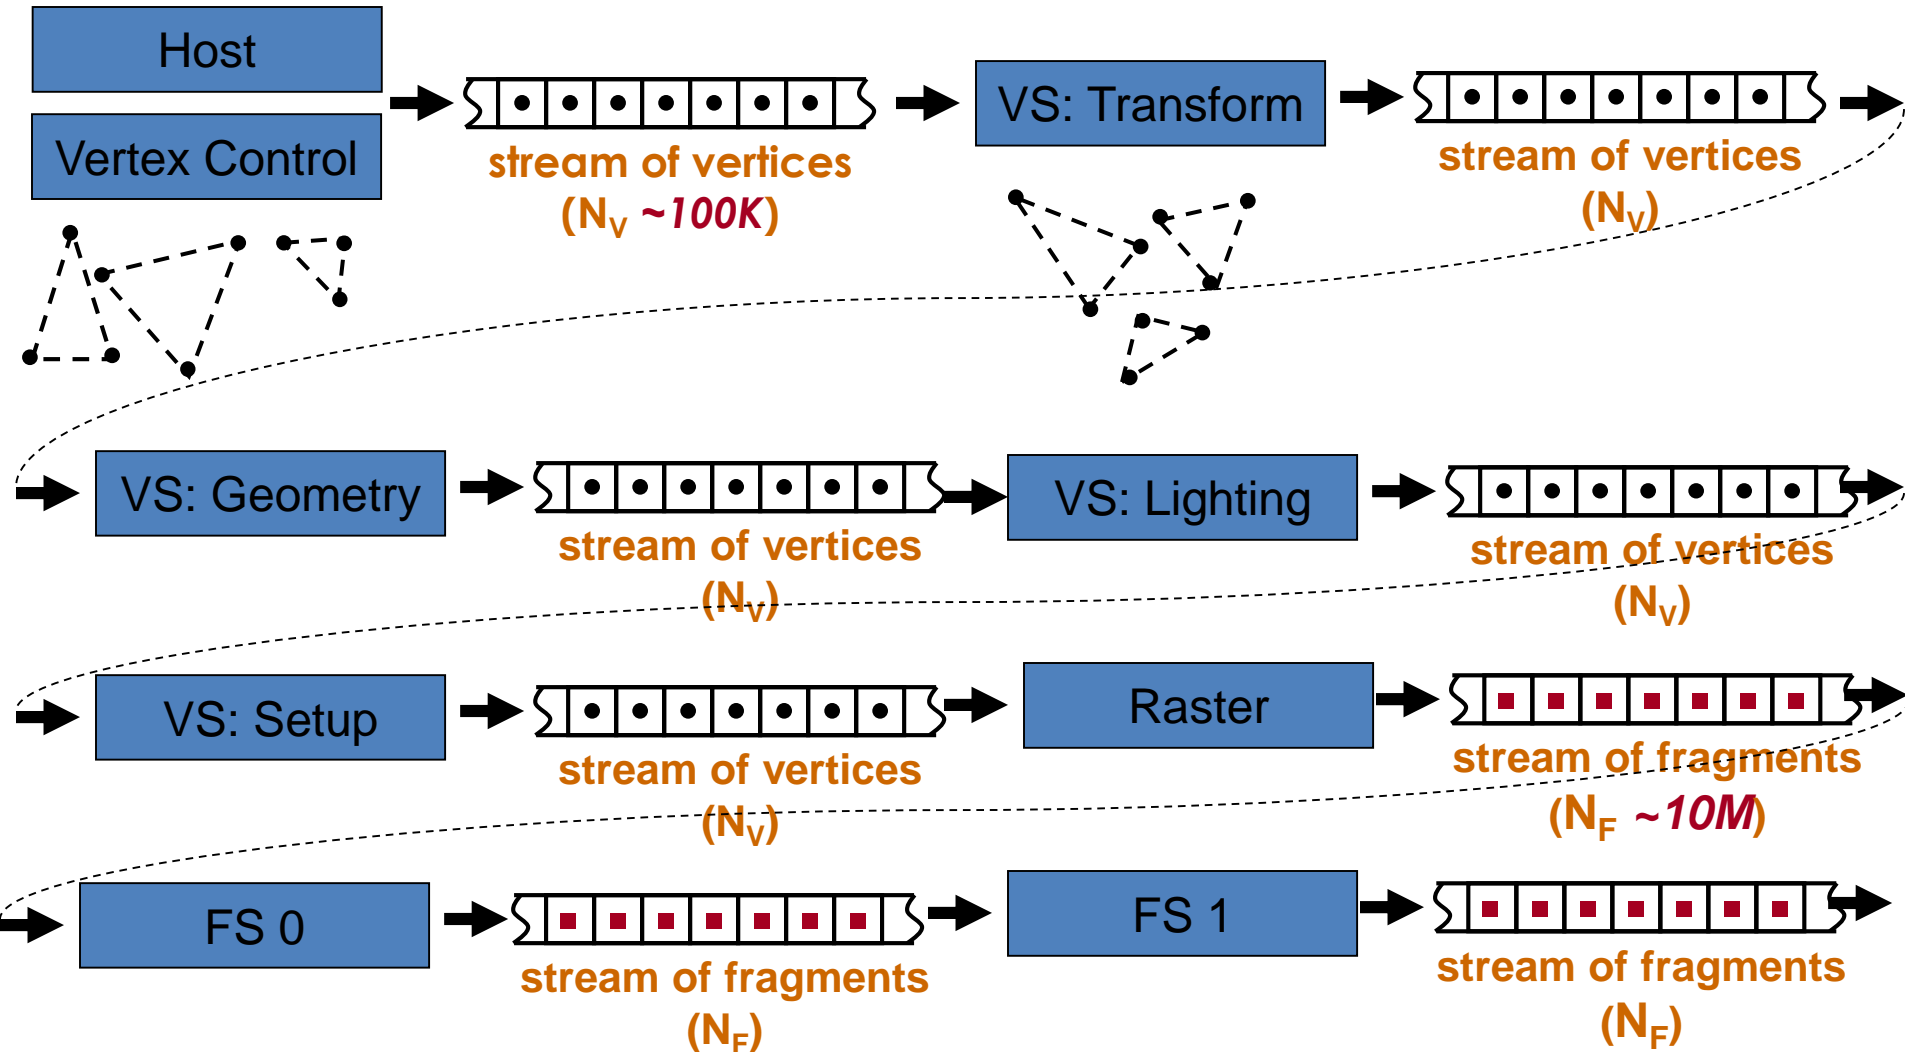


The NVIDIA GeForce Graphics Pipeline



Another View of the 3D Graphics Pipeline

8



Stream Execution Model

- Data parallel *streams* of data
- Processing *kernels*
 - Unit of Execution is processing of one stream element in one kernel – defined as a *thread*



Stream Execution Model

- Can partition the streams into chunks
 - Streams are very long and elements are independent
 - Chunks are called **strips** or **blocks**



- Unit of Execution is processing one block of data by one kernel – defined as a **thread block**

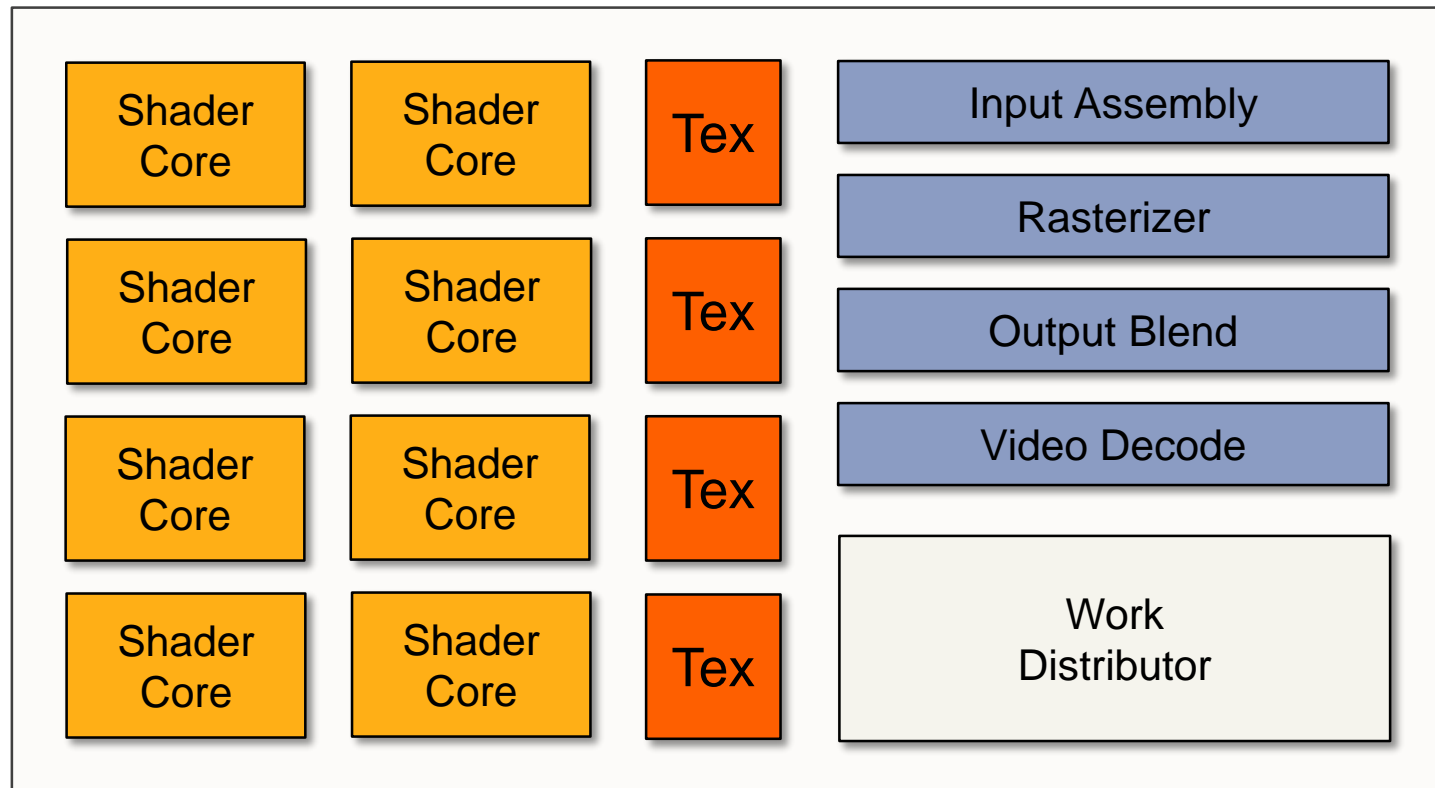


From Shader Code to a **Tera**flop: How Shader Cores Work

Kayvon Fatahalian
CMU



What's in a GPU?



Heterogeneous chip multi-processor (highly tuned for graphics)



A diffuse reflectance shader

```
sampler mySamp;  
Texture2D<float3>myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```

Independent – data parallelism



Compile shader

1 unshaded fragment input record



```
sampler mySamp;
Texture2D<float3>myTex;
float3 lightDir;

float4 diffuseShader(float3 norm, float2 uv)
{
    float3 kd;
    kd = myTex.Sample(mySamp, uv);
    kd *= clamp ( dot(lightDir, norm), 0.0, 1.0 );
    return float4(kd, 1.0);
}
```



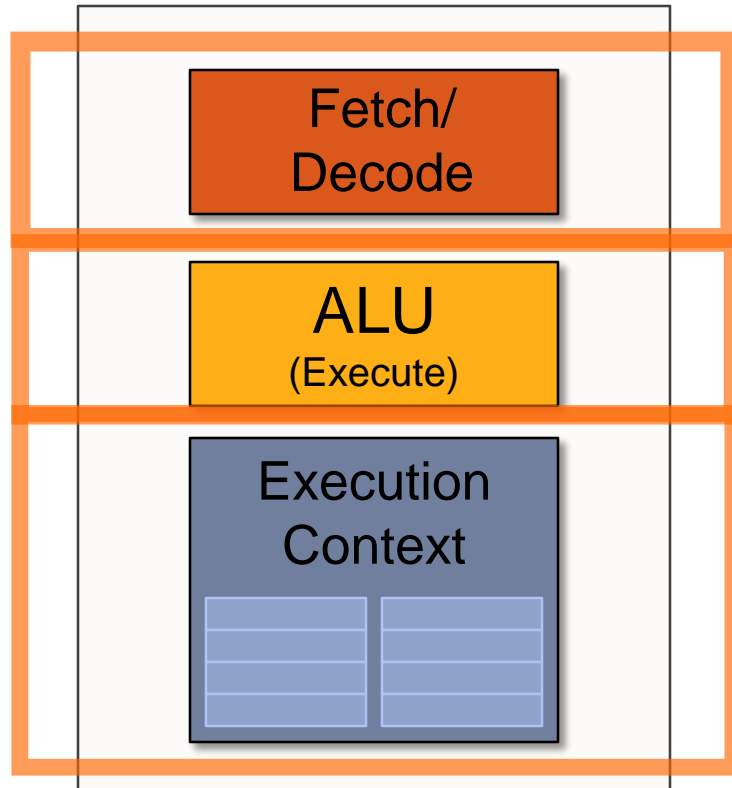
```
<diffuseShader>:
sample r0, v4, t0, s0
mul  r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, l(0.0), l(1.0)
mul  o0, r0, r3
mul  o1, r1, r3
mul  o2, r2, r3
mov  o3, l(1.0)
```



1 shaded fragment output record



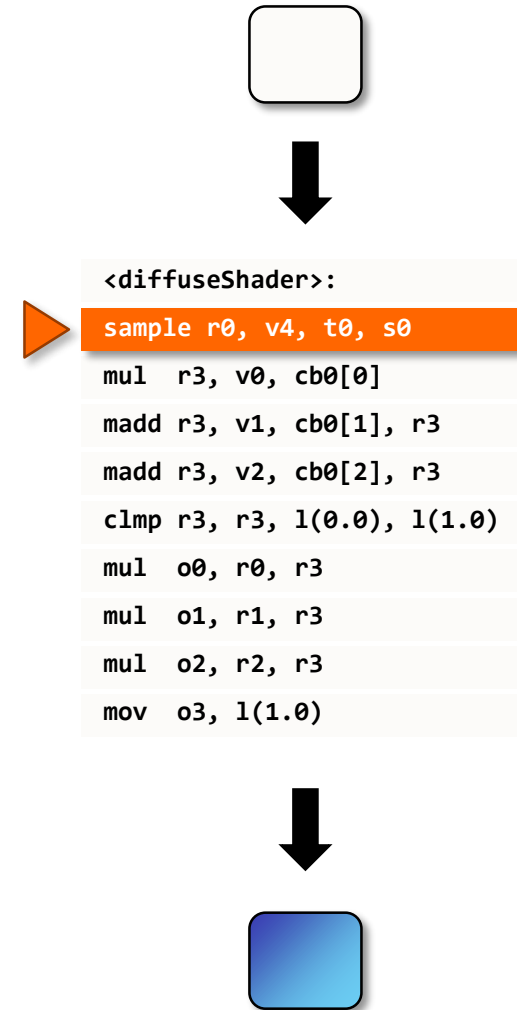

Execute shader



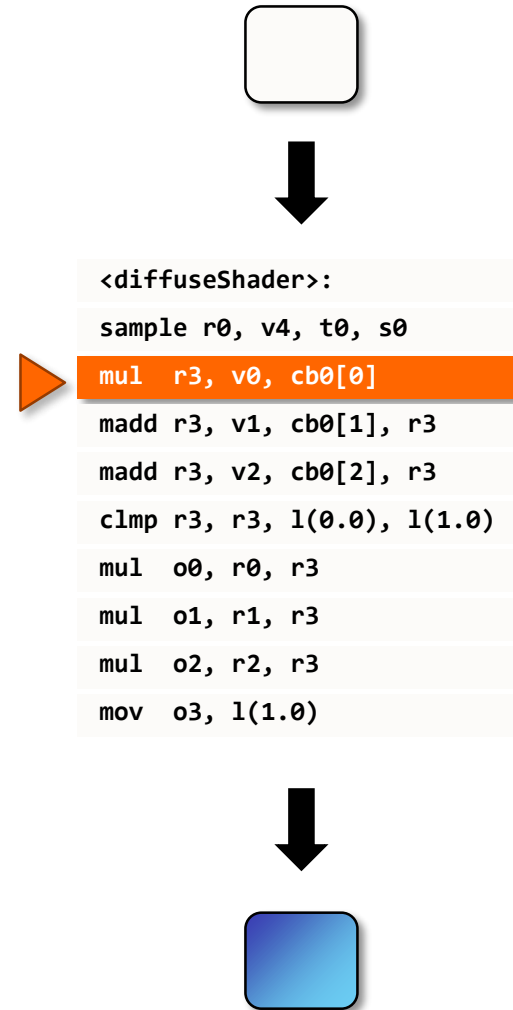

```
<diffuseShader>:
sample r0, v4, t0, s0
mul  r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, l(0.0), l(1.0)
mul  o0, r0, r3
mul  o1, r1, r3
mul  o2, r2, r3
mov  o3, l(1.0)
```



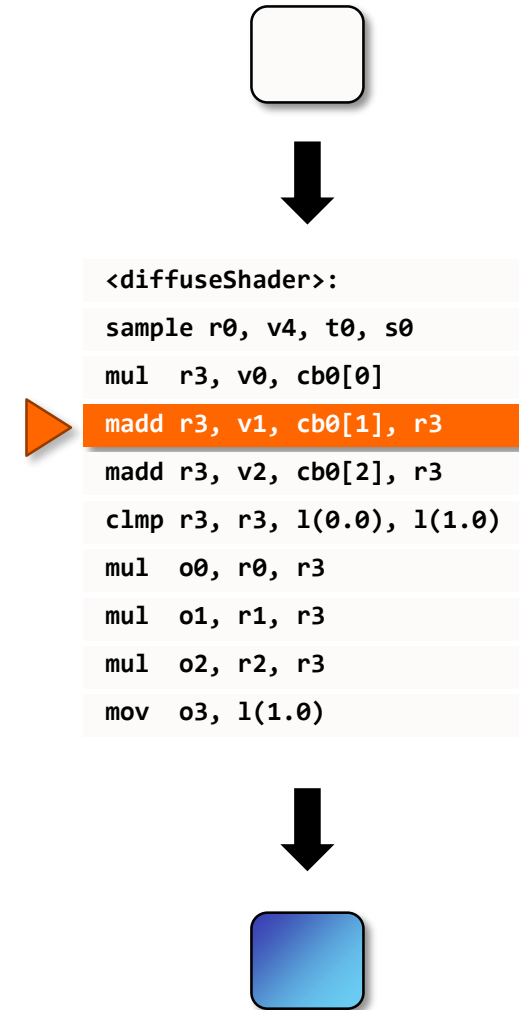

Execute shader



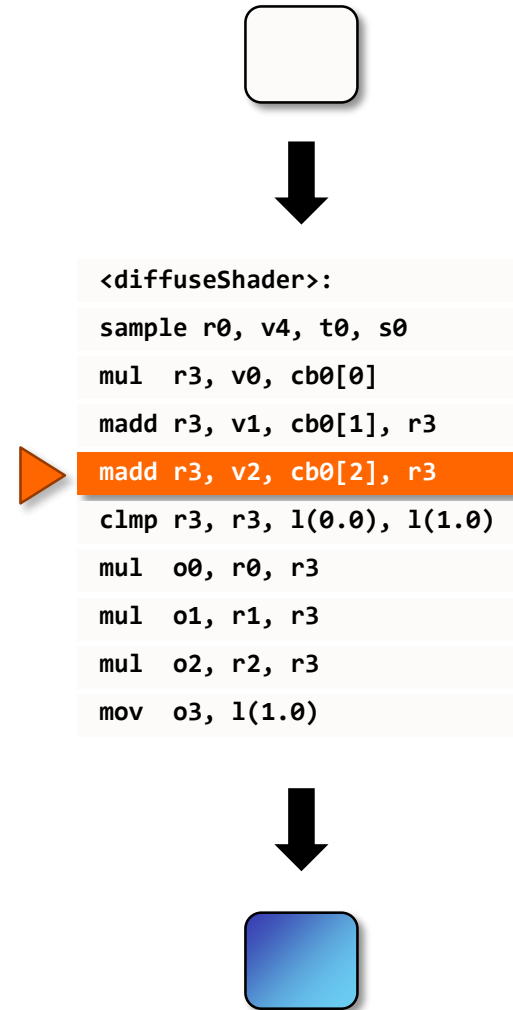

Execute shader




Execute shader



Execute shader



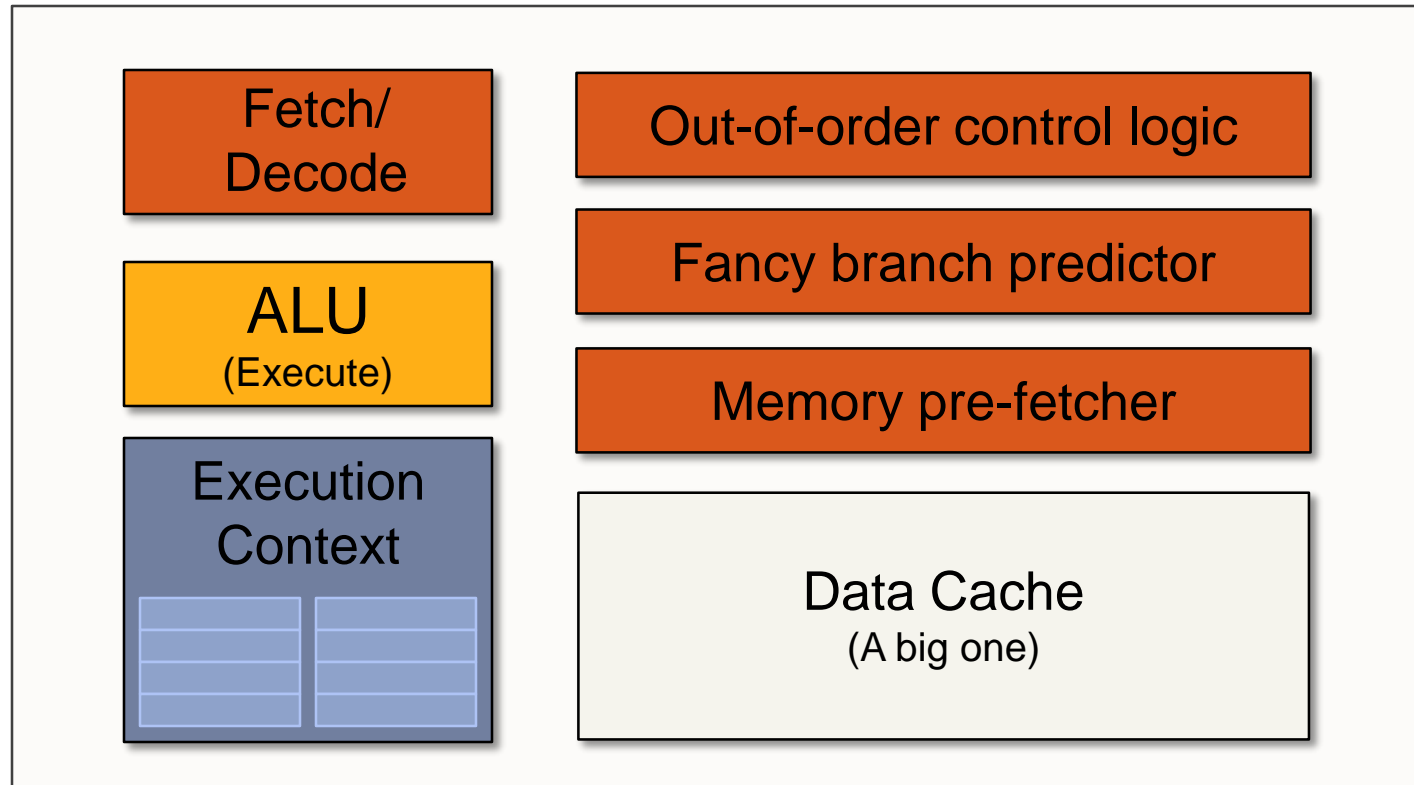
Execute shader



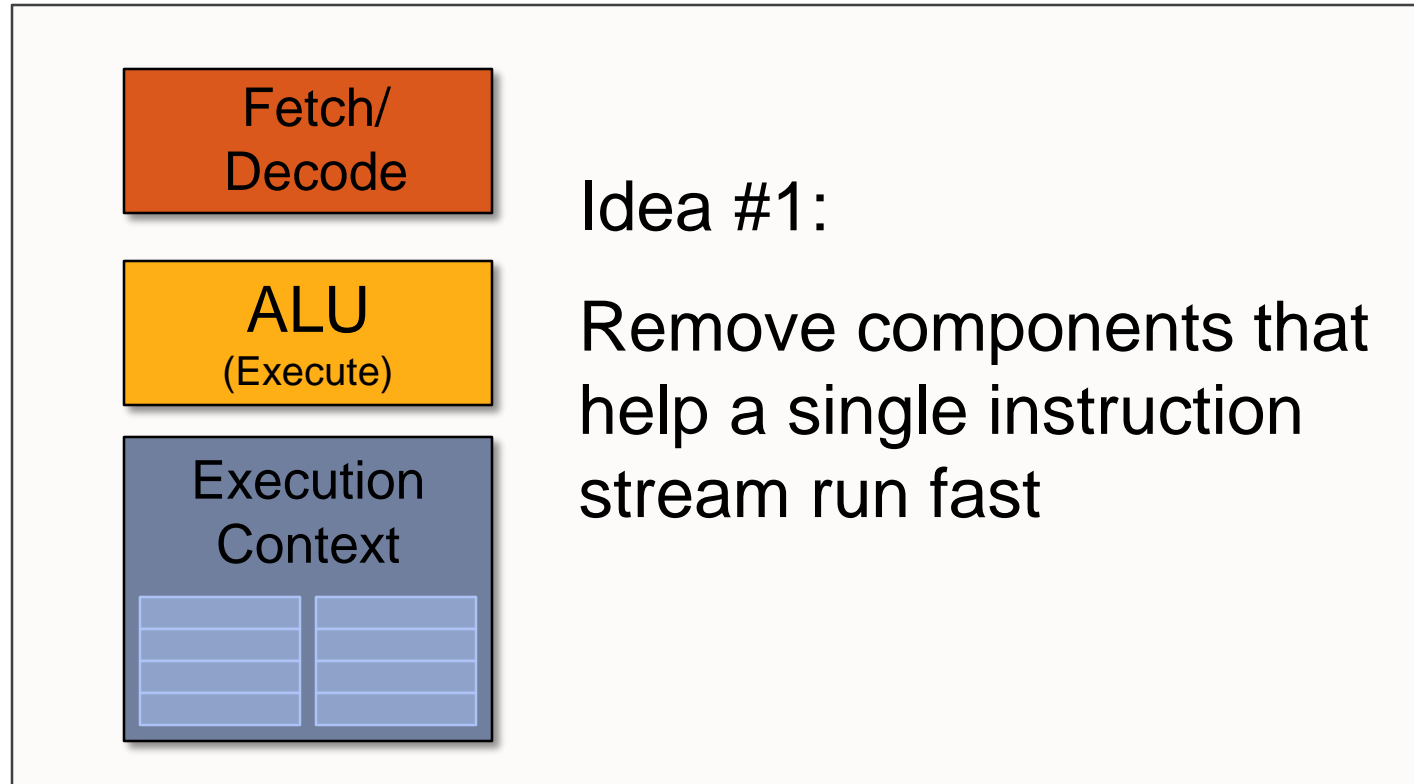
```
<diffuseShader>:
sample r0, v4, t0, s0
mul  r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, l(0.0), l(1.0)
mul  o0, r0, r3
mul  o1, r1, r3
mul  o2, r2, r3
mov  o3, l(1.0)
```



CPU-“style” cores

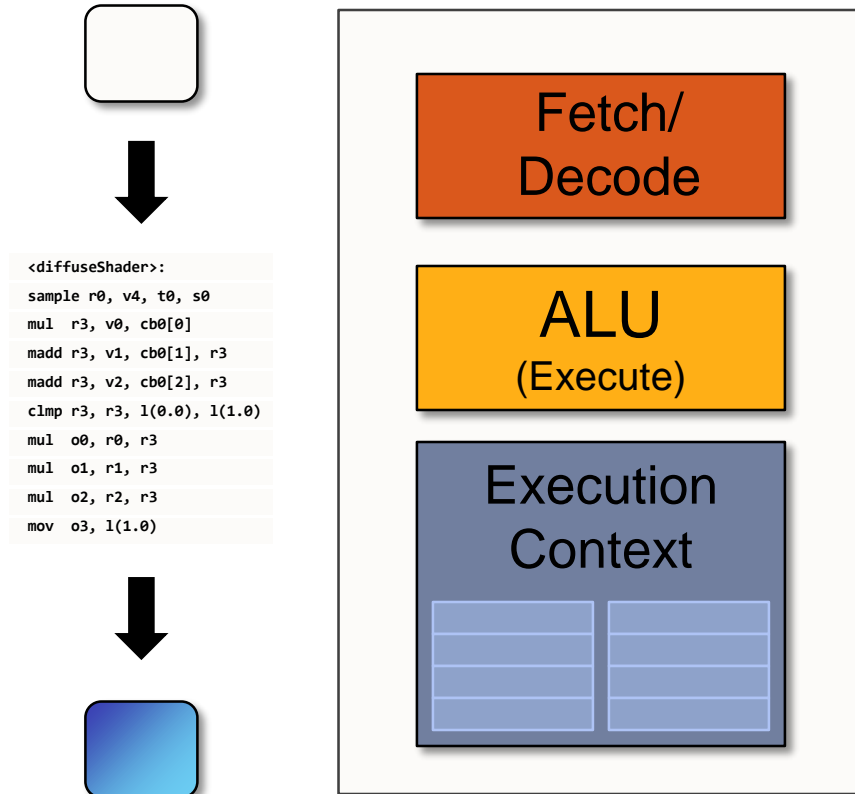


Slimming down

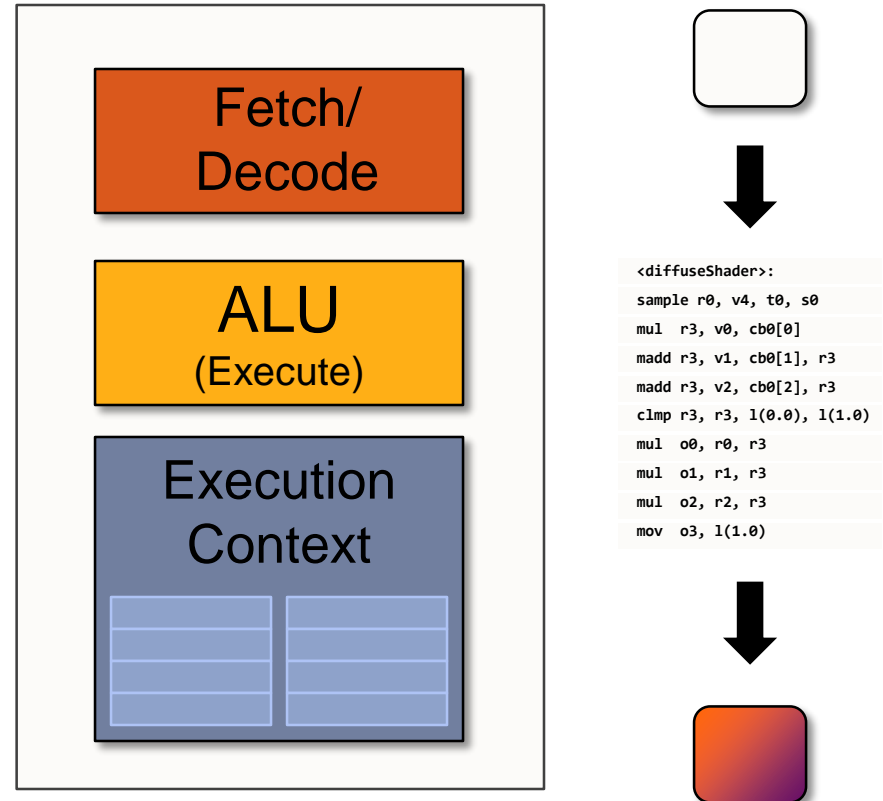


Two cores (two fragments in parallel)






fragment 1




fragment 2



Four cores (four fragments in parallel)




Sixteen cores (sixteen fragments in parallel)



16 cores = 16 simultaneous instruction streams



Instruction stream coherence

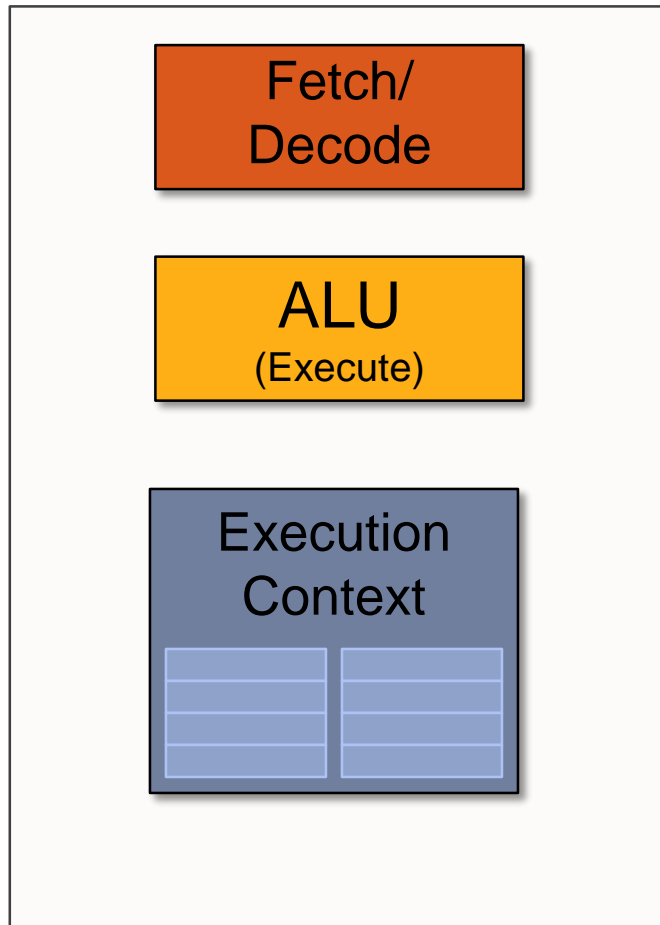


But... many fragments should be able to share an instruction stream!

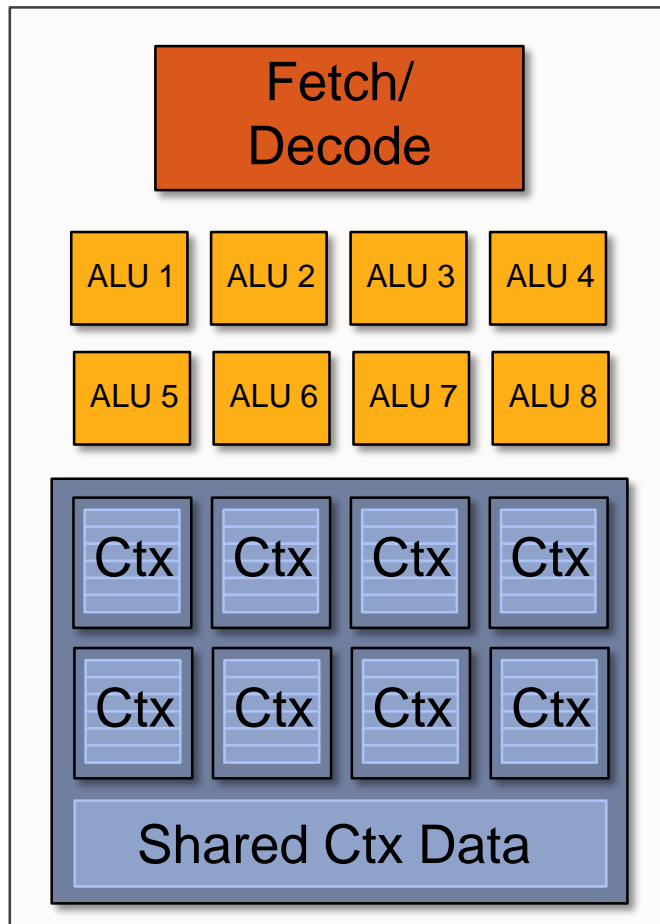
```
<diffuseShader>:
sample r0, v4, t0, s0
mul  r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, l(0.0), l(1.0)
mul  o0, r0, r3
mul  o1, r1, r3
mul  o2, r2, r3
mov  o3, l(1.0)
```



Recall: simple processing core



Add ALUs



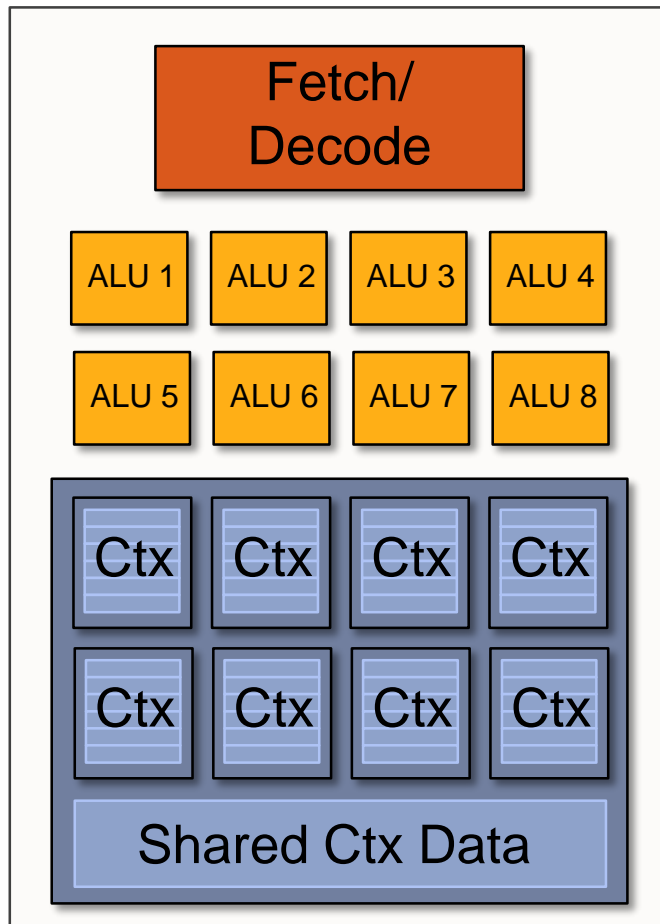
Idea #2:

Amortize cost/complexity of managing an instruction stream across many ALUs

SIMD processing



Modifying the shader




```
<diffuseShader>:
sample r0, v4, t0, s0
mul  r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, l(0.0), l(1.0)
mul  o0, r0, r3
mul  o1, r1, r3
mul  o2, r2, r3
mov  o3, l(1.0)
```

Original compiled shader:
Processes one fragment
using scalar ops on scalar
registers



Modifying the shader




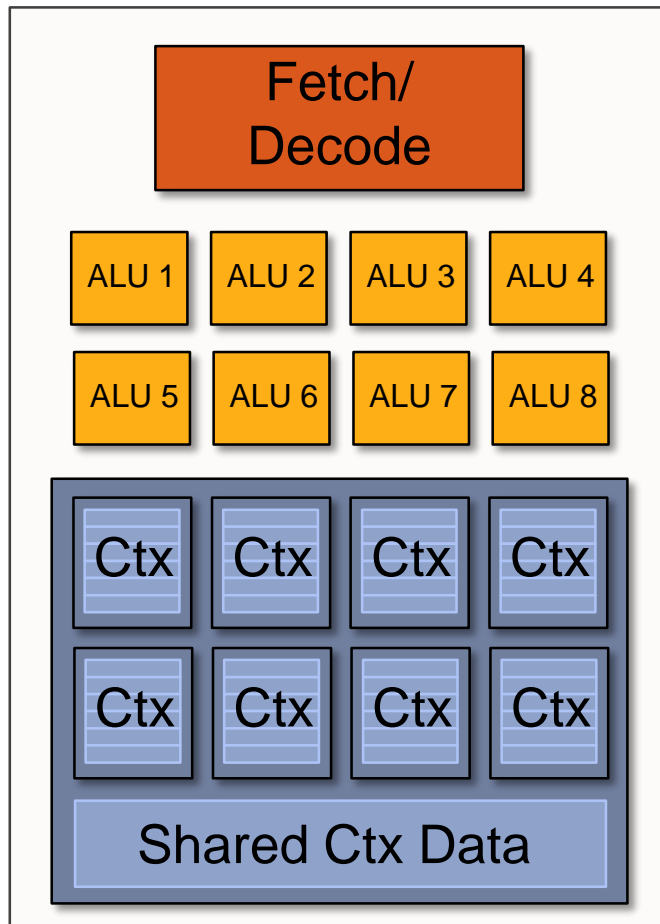
```
<VEC8_diffuseShader>:
VEC8_sample vec_r0, vec_v4, t0, vec_s0
VEC8_mul  vec_r3, vec_v0, cb0[0]
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3
VEC8_clmp vec_r3, vec_r3, 1(0.0), 1(1.0)
VEC8_mul  vec_o0, vec_r0, vec_r3
VEC8_mul  vec_o1, vec_r1, vec_r3
VEC8_mul  vec_o2, vec_r2, vec_r3
VEC8_mov  vec_o3, 1(1.0)
```

New compiled shader:

Processes 8 fragments
using vector ops on vector
registers



Modifying the shader





<VEC8_diffuseShader>:

```
VEC8_sample vec_r0, vec_v4, t0, vec_s0
VEC8_mul  vec_r3, vec_v0, cb0[0]
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3
VEC8_clmp vec_r3, vec_r3, l(0.0), l(1.0)
VEC8_mul  vec_o0, vec_r0, vec_r3
VEC8_mul  vec_o1, vec_r1, vec_r3
VEC8_mul  vec_o2, vec_r2, vec_r3
VEC8_mov  vec_o3, l(1.0)
```





128 fragments in parallel




16 cores = 128 ALUs
 = 16 simultaneous instruction streams



128 [vertices / fragments primitives CUDA threads OpenCL work items compute shader threads] in parallel



But what about branches?




<unconditional
shader code>

```
if (x > 0) {
    y = pow(x, exp);
    y *= Ks;
    refl = y + Ka;
} else {
    x = 0;
    refl = Ka;
}
```

<resume unconditional
shader code>



But what about branches?



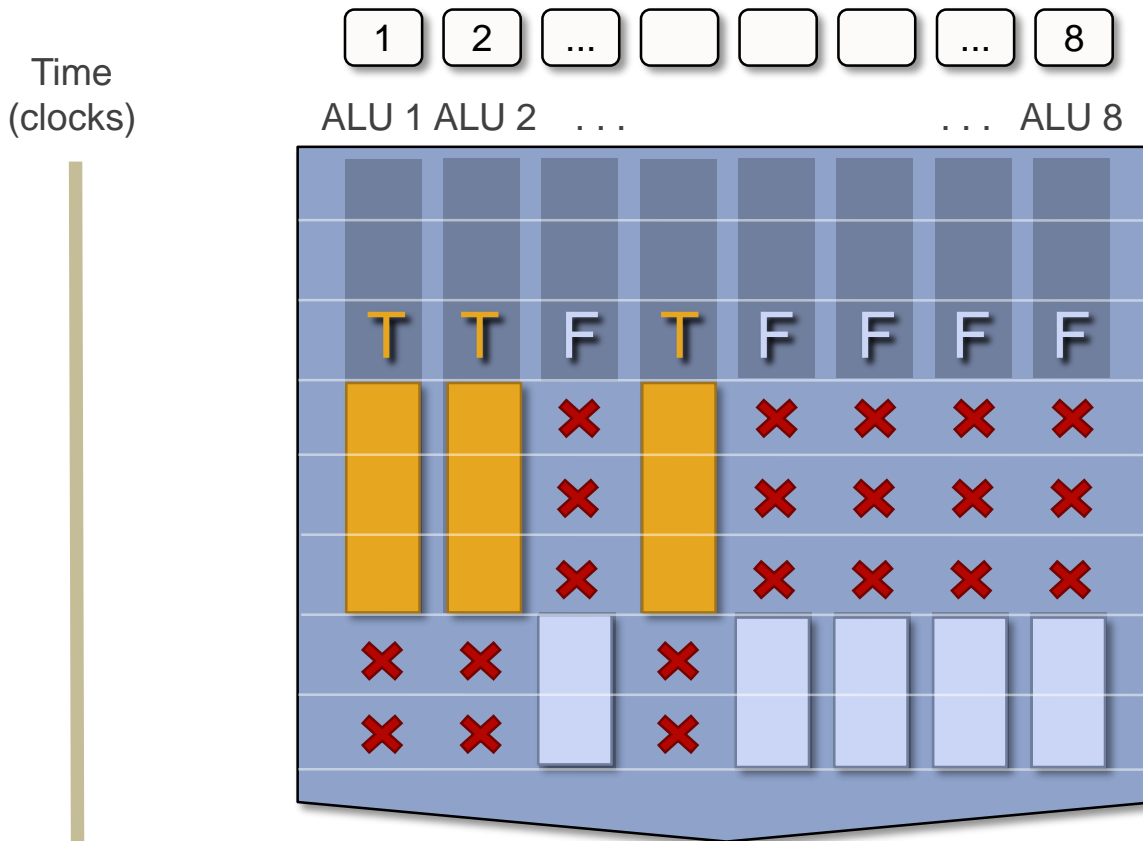
<unconditional
shader code>

```
if (x > 0) {
    y = pow(x, exp);
    y *= Ks;
    refl = y + Ka;
} else {
    x = 0;
    refl = Ka;
}
```

<resume unconditional
shader code>



But what about branches?



Not all ALUs do useful work!
Worst case: 1/8
performance

<unconditional
shader code>

```
if (x > 0) {
```

```
    y = pow(x, exp);
```

```
    y *= Ks;
```

```
    refl = y + Ka;
```

```
} else {
```

```
    x = 0;
```


```
    refl = Ka;
```

```
}
```

<resume unconditional
shader code>



But what about branches?



<unconditional
shader code>

```
if (x > 0) {
```

```
    y = pow(x, exp);
```

```
    y *= Ks;
```

```
    refl = y + Ka;
```

```
} else {
```

```
    x = 0;
```

```
    refl = Ka;
```

```
}
```


<resume unconditional
shader code>



Clarification

SIMD processing does not imply SIMD instructions

- Option 1: Explicit vector instructions
 - Intel/AMD x86 SSE, Intel Larrabee
- Option 2: Scalar instructions, implicit HW vectorization
 - HW determines instruction stream sharing across ALUs (amount of sharing hidden from software)
 - NVIDIA GeForce (“SIMT” warps), ATIRadeon architectures



In practice: 16 to 64 fragments share an instruction stream



Stalls!

Stalls occur when a core cannot run the next instruction because of a dependency on a previous operation.

Texture access latency = 100's to 1000's of cycles

We've removed the fancy caches and logic that helps avoid stalls.



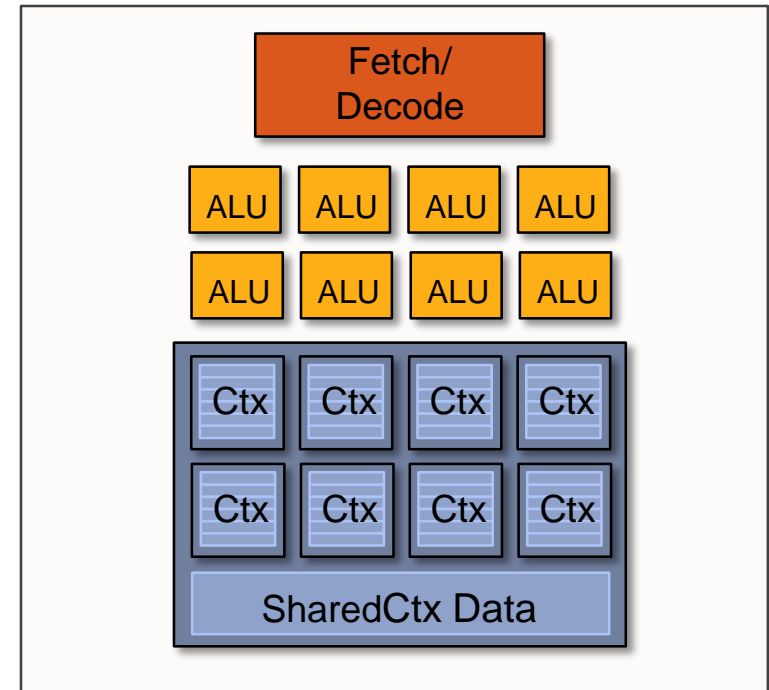

But we have **LOTS** of independent fragments.

Idea #3:


Interleave processing of many fragments on a single core to avoid stalls caused by high latency operations.




Hiding shader stalls



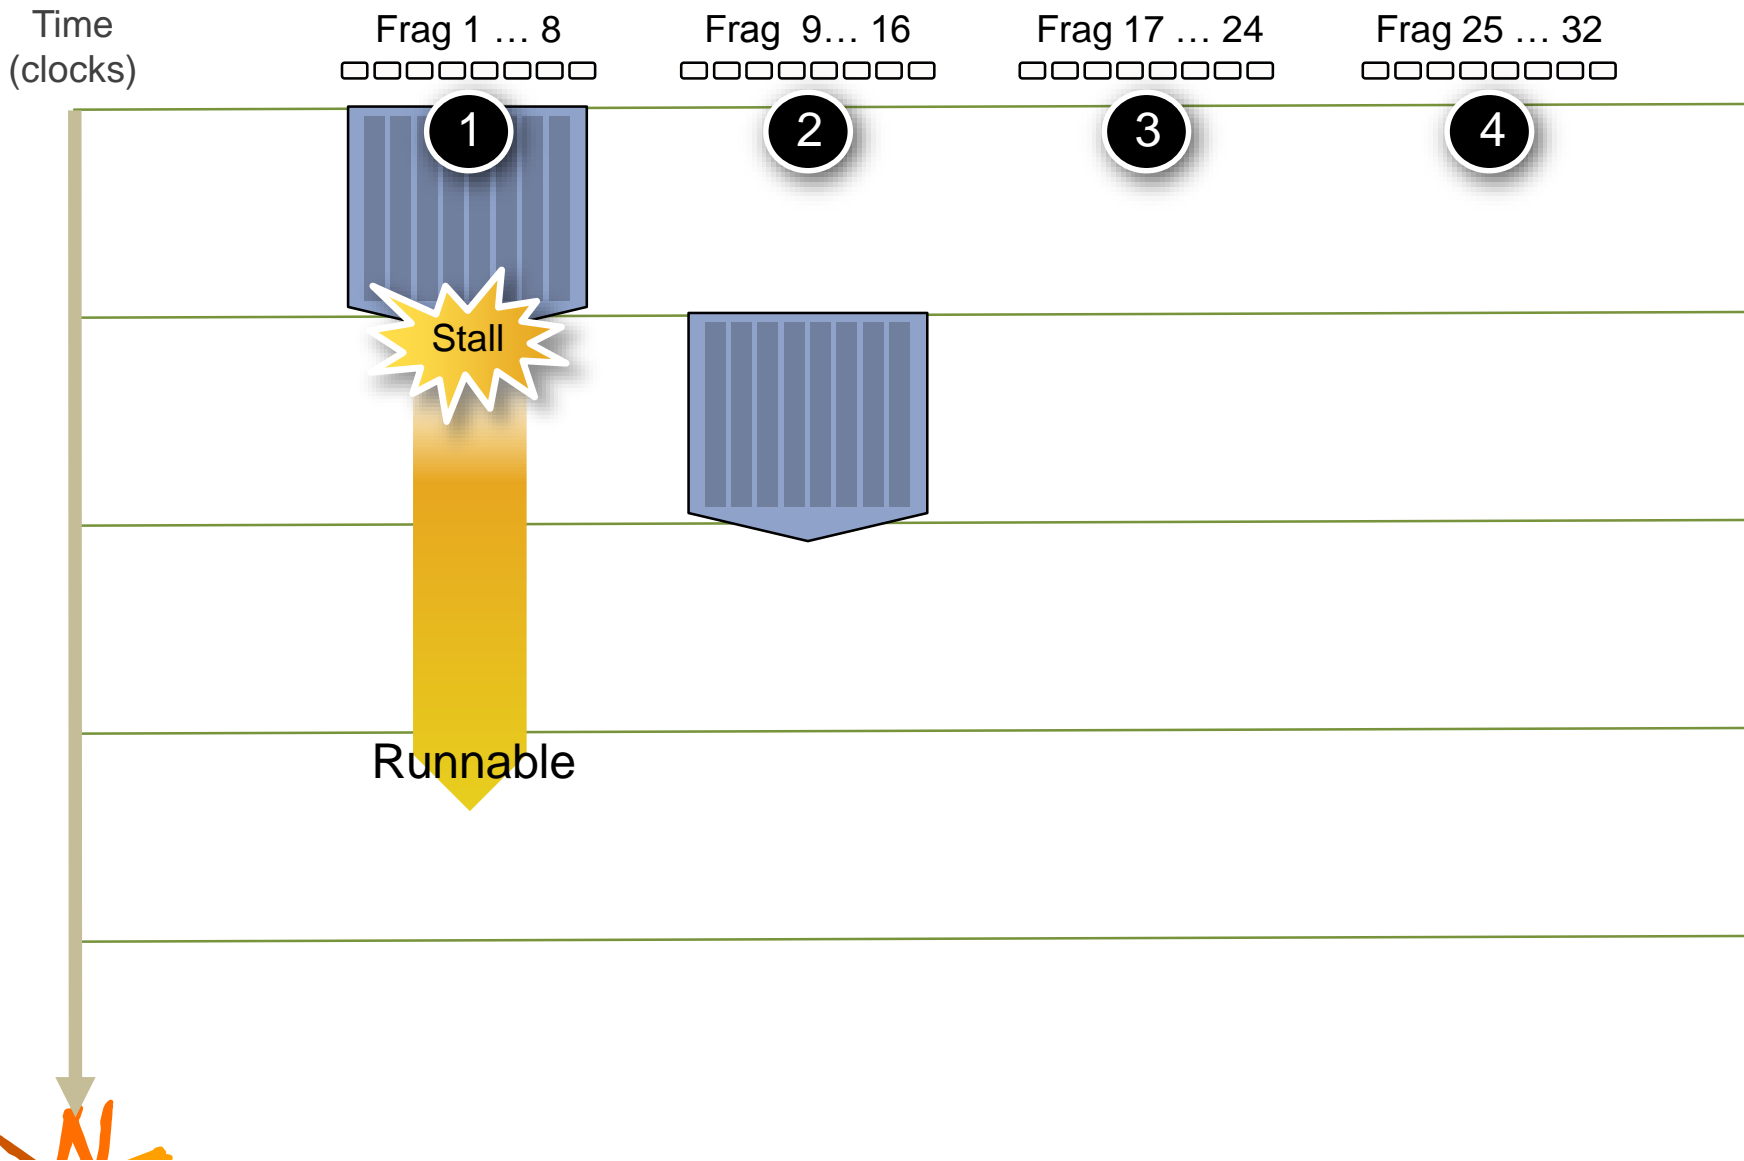
Hiding shader stalls



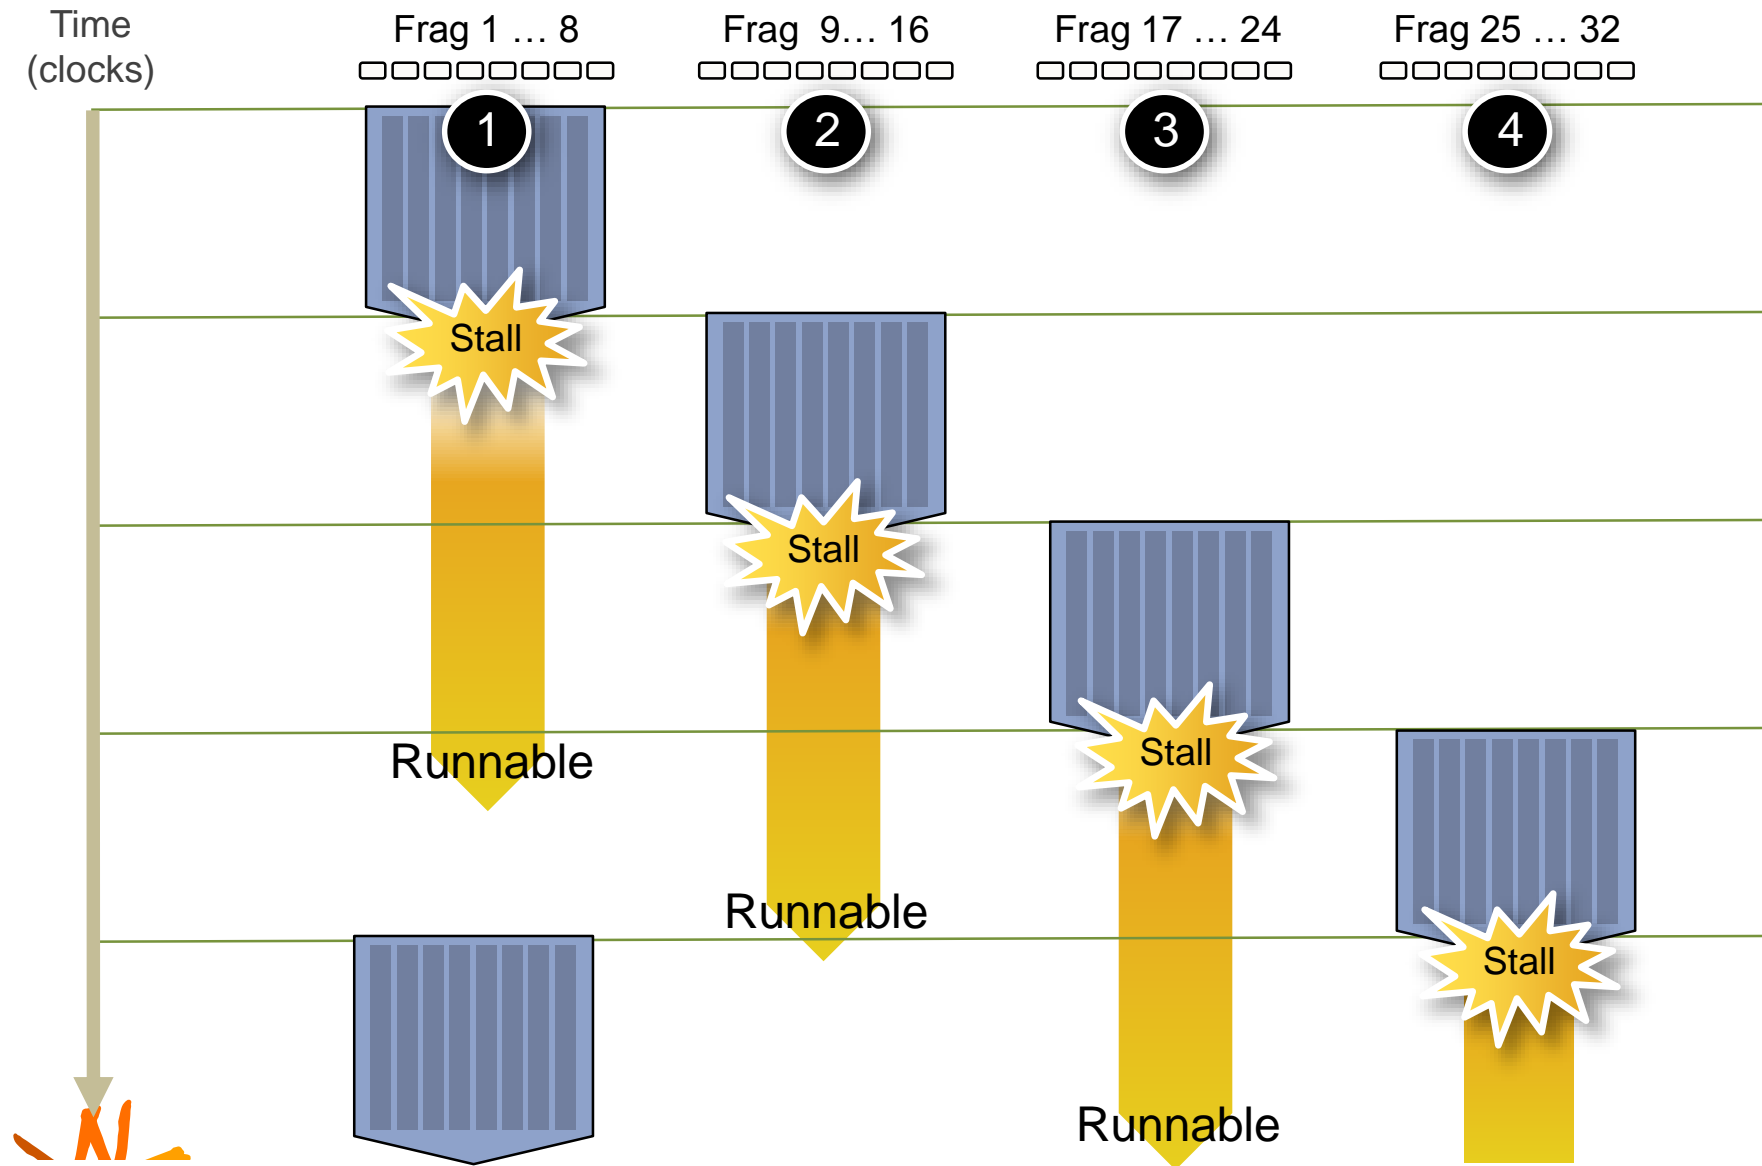
Hiding shader stalls



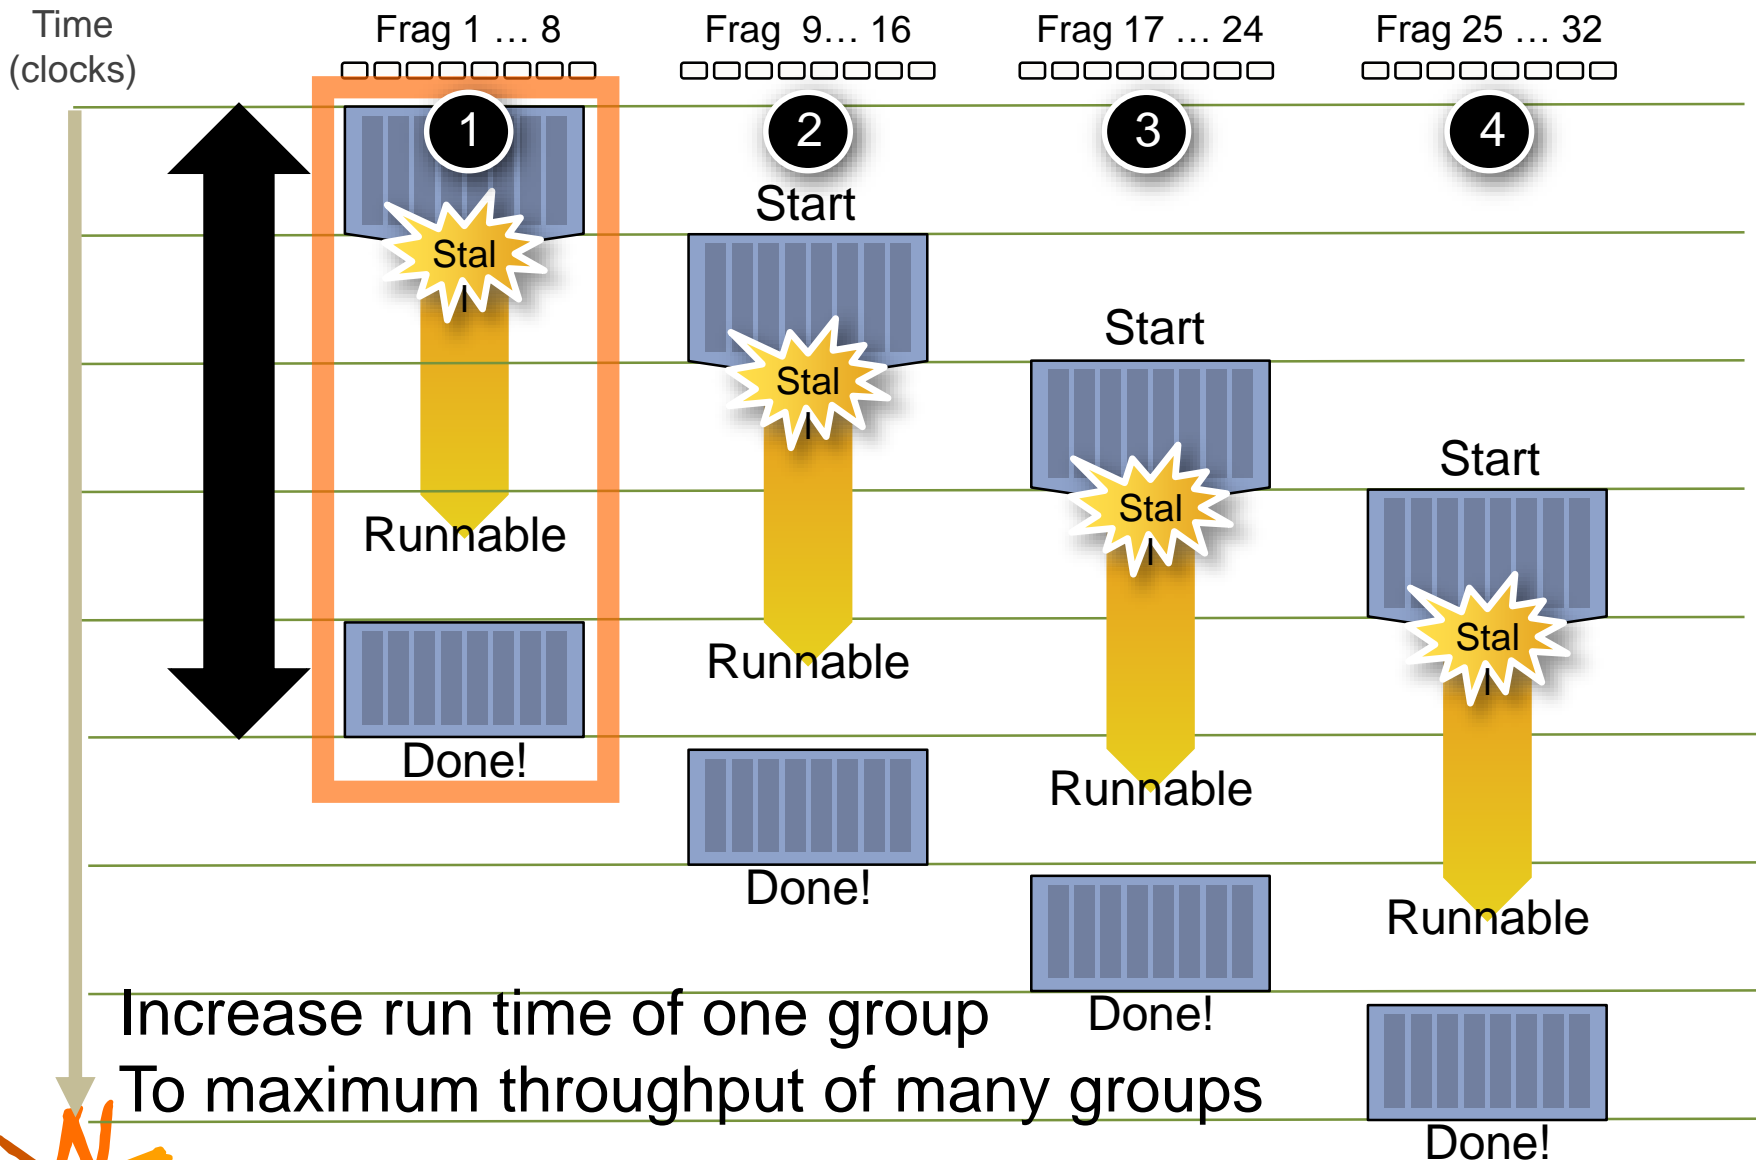
Hiding shader stalls



Hiding shader stalls




Throughput!



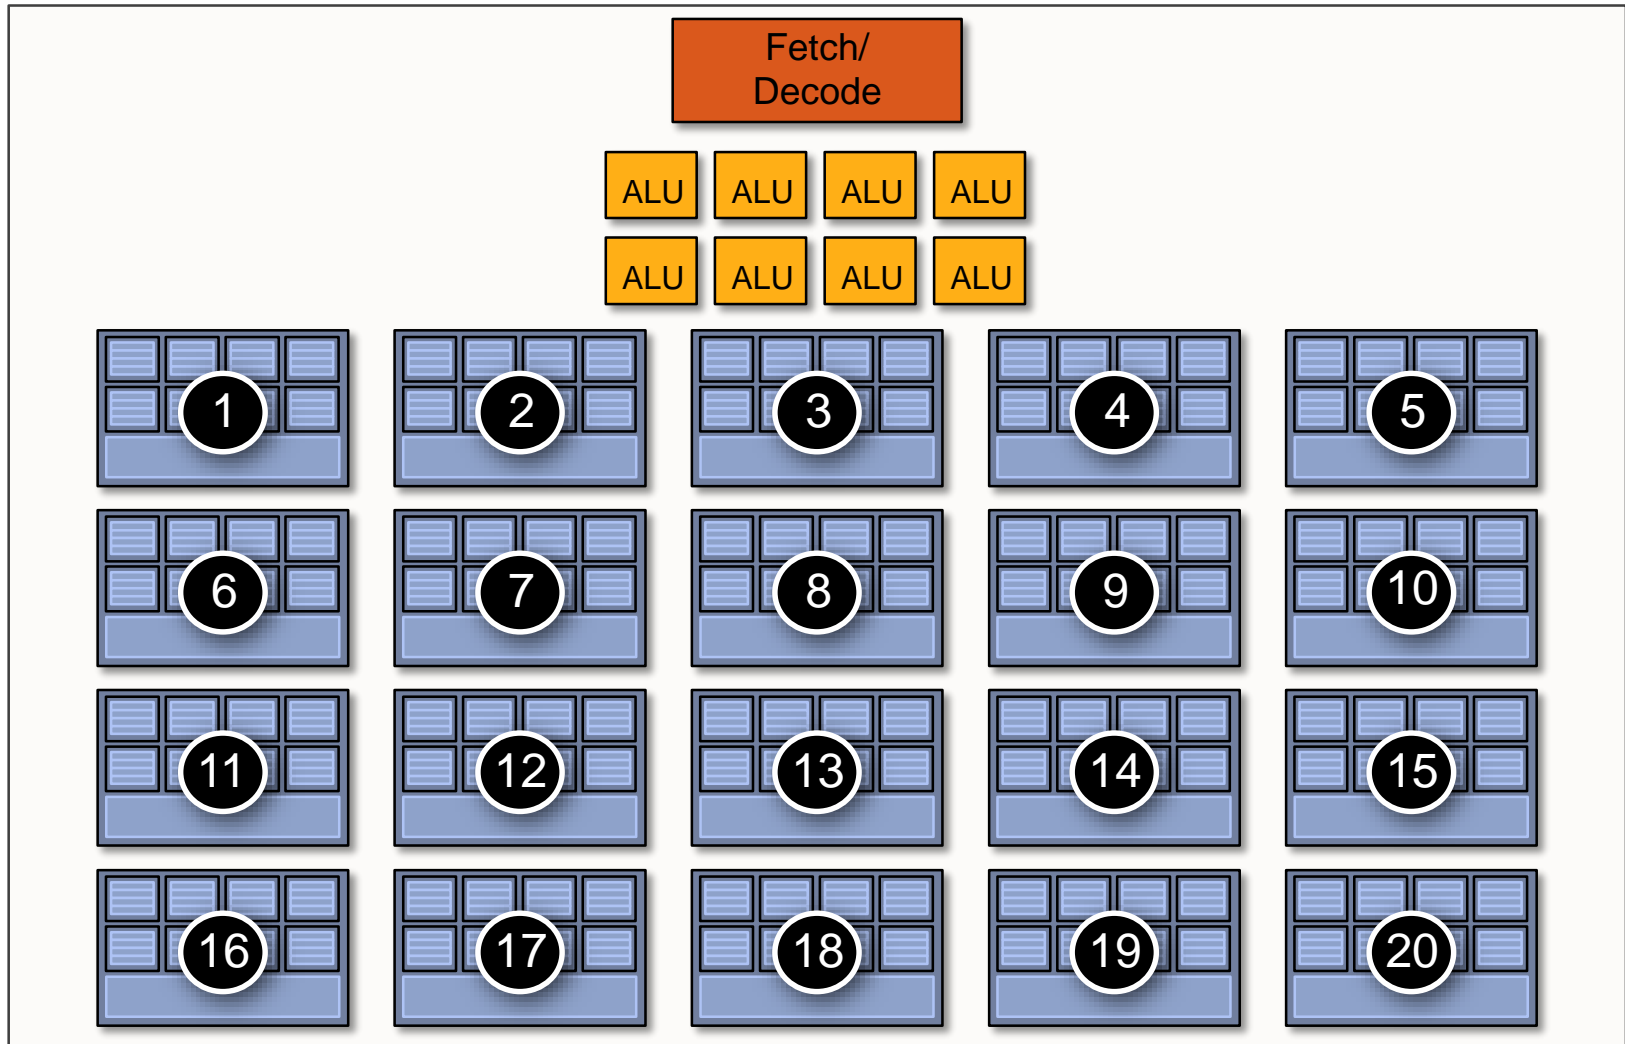
Storing contexts

47

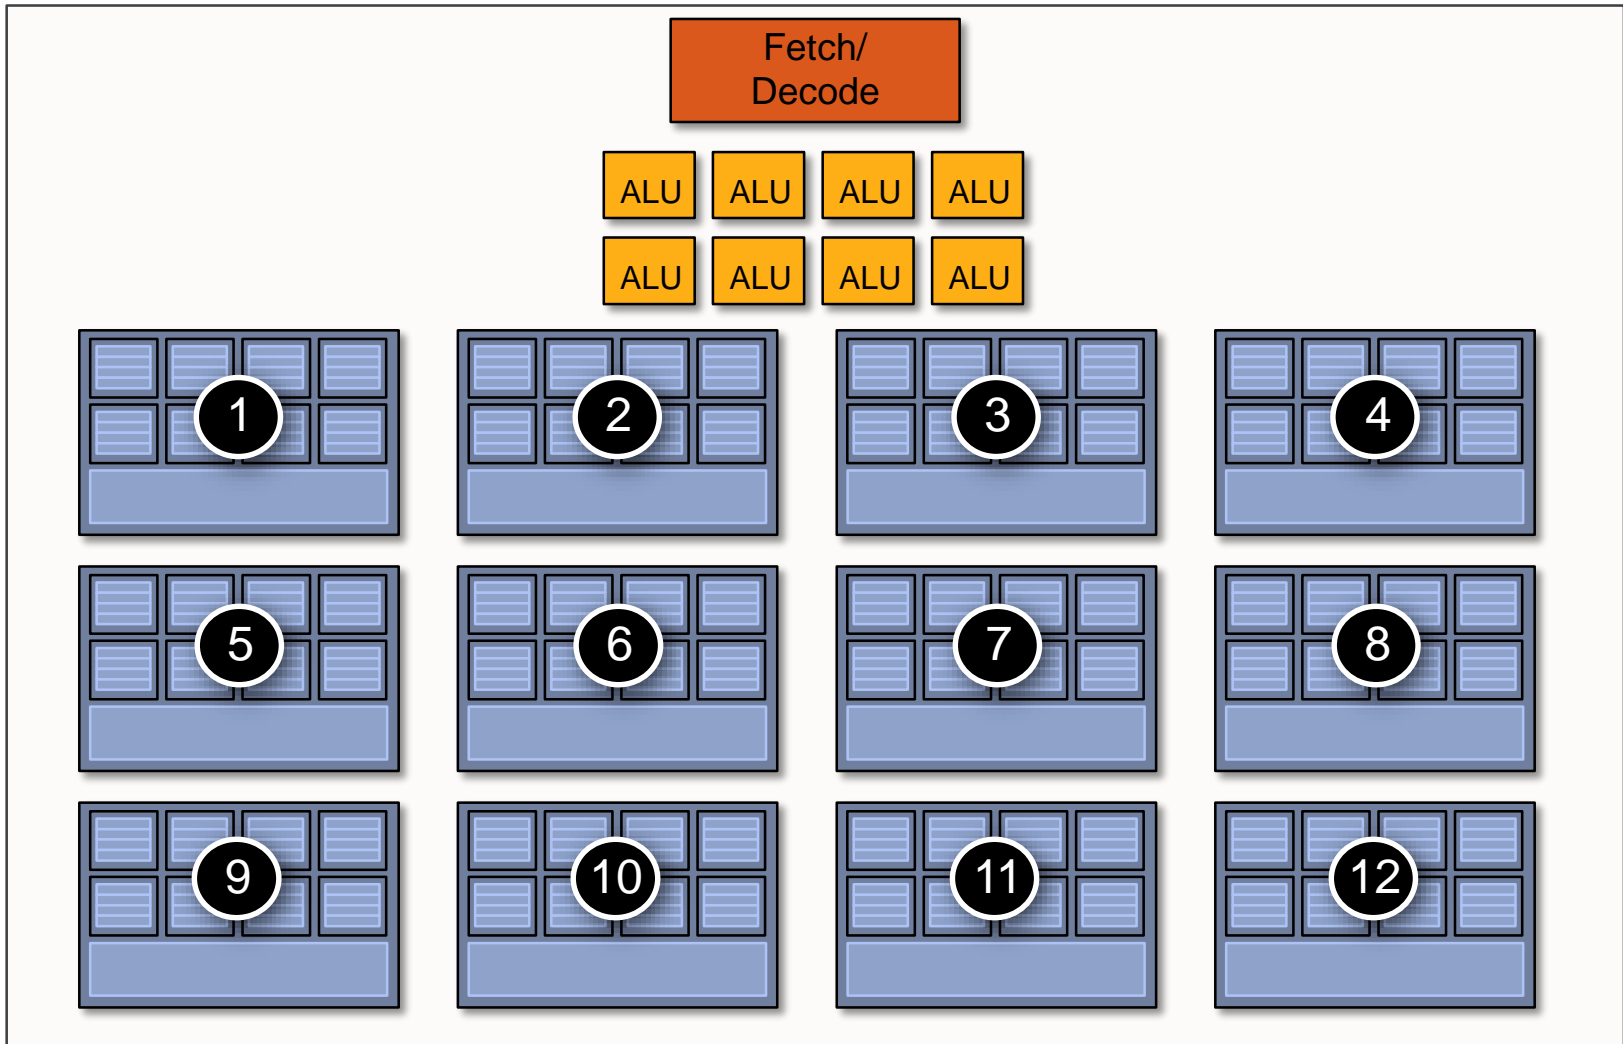


Twenty small contexts

(maximal latency hiding ability)




Twelve medium contexts



Four large contexts

(low latency hiding ability)



Summary: three key ideas for GPU architecture


1. Use many “slimmed down cores” to run in parallel
2. Pack cores full of ALUs (by sharing instruction stream across groups of fragments)
 - Option 1: Explicit SIMD vector instructions
 - Option 2: Implicit sharing managed by hardware
3. Avoid latency stalls by interleaving execution of many groups of fragments
 - When one group stalls, work on another group



GPU block diagram key



= single “physical” instruction stream fetch/decode
(functional unit control)



= SIMD programmable functional unit (FU), control shared with other
functional units. This functional unit may contain multiple 32-bit “ALUs”



= 32-bit mul-add unit



= 32-bit multiply unit



= execution context storage



= fixed function unit



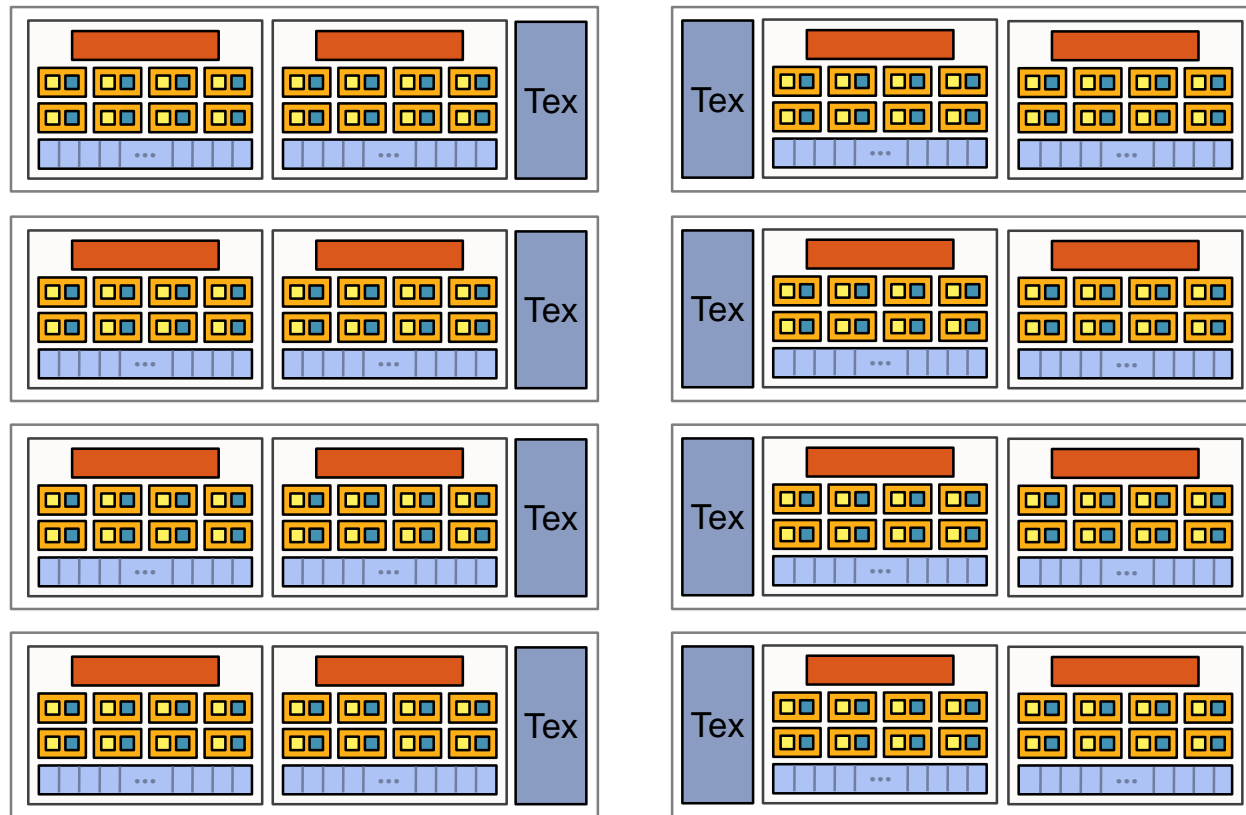
NVIDIA GeForce 8800/9800

- NVIDIA-speak:
 - 128 stream processors
 - “SIMT execution” (automatic HW-managed sharing of instruction stream)
- Generic speak:
 - 16 processing cores
 - 8 SIMD functional units per core
 - 1 mul-add (2 flops) + 1 mul per functional units (3 flops/clock)
 - Best case: 128 mul-adds + 128 muls per clock
 - 1.3 GHz clock
 - $16 * 8 * (2 + 1) * 1.2 = 460$ GFLOPS
- Mapping data-parallelism to chip:
 - Instruction stream shared across 32 fragments (16 for vertices)
 - 8 fragments run on 8 SIMD functional units in one clock
 - Instruction repeated for 4 clocks (2 clocks for vertices)



NVIDIA GeForce 8800/9800

55



Zcull/Clip/Rast

Output Blend

Work Distributor



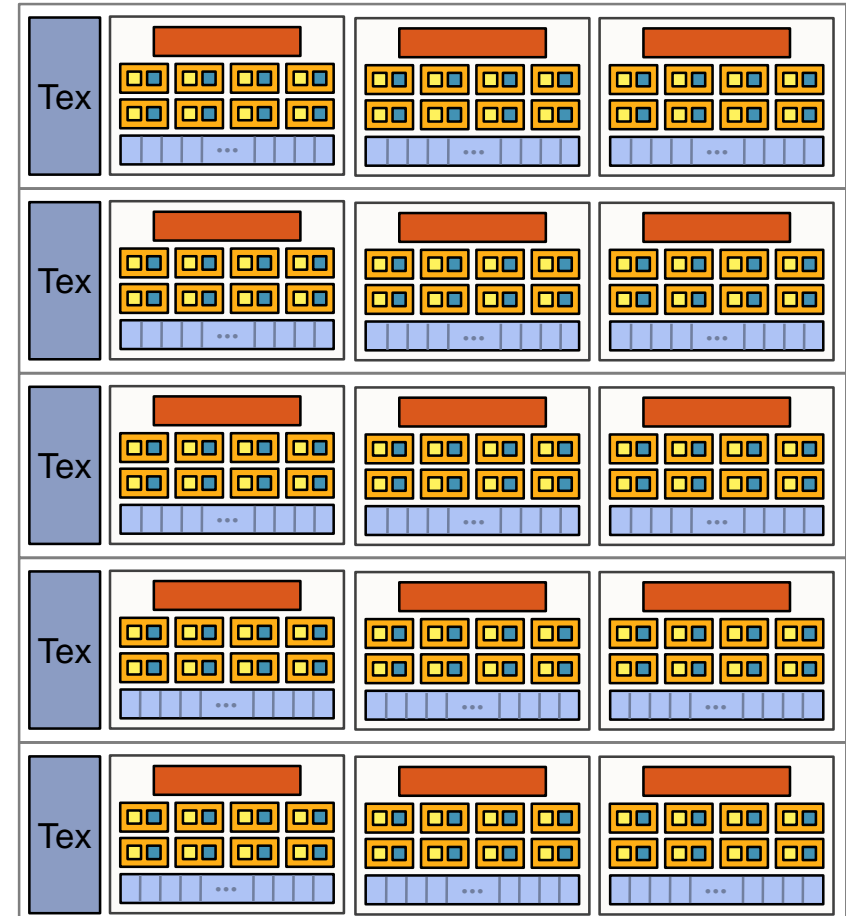
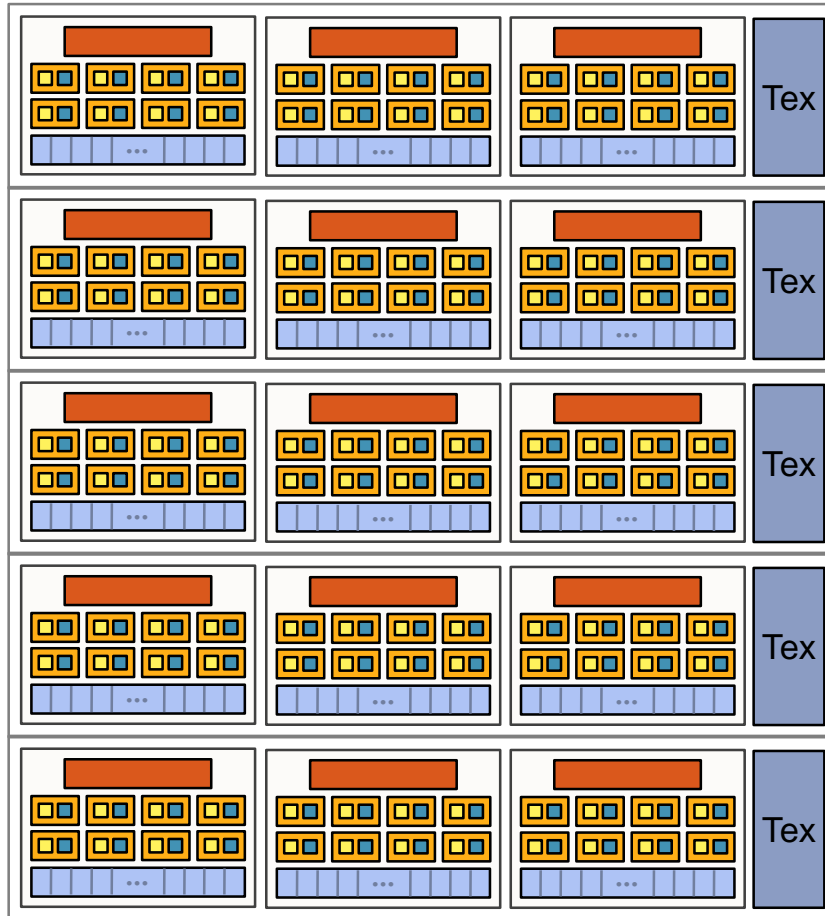
NVIDIA GeForce GTX 280

- NVIDIA-speak:
 - 240 stream processors
 - “SIMT execution” (automatic HW-managed sharing of instruction stream)
- Generic speak:
 - 30 processing cores
 - 8 SIMD functional units per core
 - 1 mul-add (2 flops) + 1 mul per functional units (3 flops/clock)
 - Best case: 240 mul-adds + 240 muls per clock
 - 1.3 GHz clock
 - $30 * 8 * (2 + 1) * 1.3 = 933$ GFLOPS
- Mapping data-parallelism to chip:
 - Instruction stream shared across 32 “threads”
 - 16 threads run on 8 SIMD functional units in one clock
 - Instruction repeated for 4 clocks (2 clocks for vertices)



NVIDIA GeForce GTX 280

57



Zcull/Clip/Rast

Output Blend

Work Distributor



NVIDIA Fermi

- NVIDIA-speak:
 - 512 CUDA processors (formerly stream processors)
 - “SIMT execution” (automatic HW-managed sharing of instruction stream)
- Generic speak:
 - 26 processing cores
 - 32 SIMD functional units per core
 - 1 FPU (FMA) + 1 INT per functional units (2 flops/clock)
 - Best case: 512 FP mul-adds
 - 1.3 GHz clock
 - $16 * 32 * 2 * 1.5 = 1.5$ GFLOPS (more real than earlier)
- Mapping data-parallelism to chip:
 - Instruction stream shared across 32 threads
 - 32 fragments run on 16 SIMD functional units in one clock
 - Instruction repeated for 2 clocks ?





ATI Radeon 4870

- AMD/ATI-speak:
 - 800 stream processors
 - Automatic HW-managed sharing of scalar instruction stream (like “SIMT”)
- Generic speak:
 - 10 processing cores
 - 16 SIMD functional units per core
 - 5 mul-adds per functional unit ($5 * 2 = 10$ flops/clock)
 - Best case: 800 mul-adds per clock
 - 750 MHz clock
 - $10 * 16 * 5 * 2 * .75 = 1.2$ TFLOPS
- Mapping data-parallelism to chip:
 - Instruction stream shared across 64 fragments
 - 16 fragments run on 16 SIMD functional units in one clock
 - Instruction repeated for 4 consecutive clocks



ATI Radeon 4870

60



Zcull/Clip/Rast

Output Blend

Work Distributor



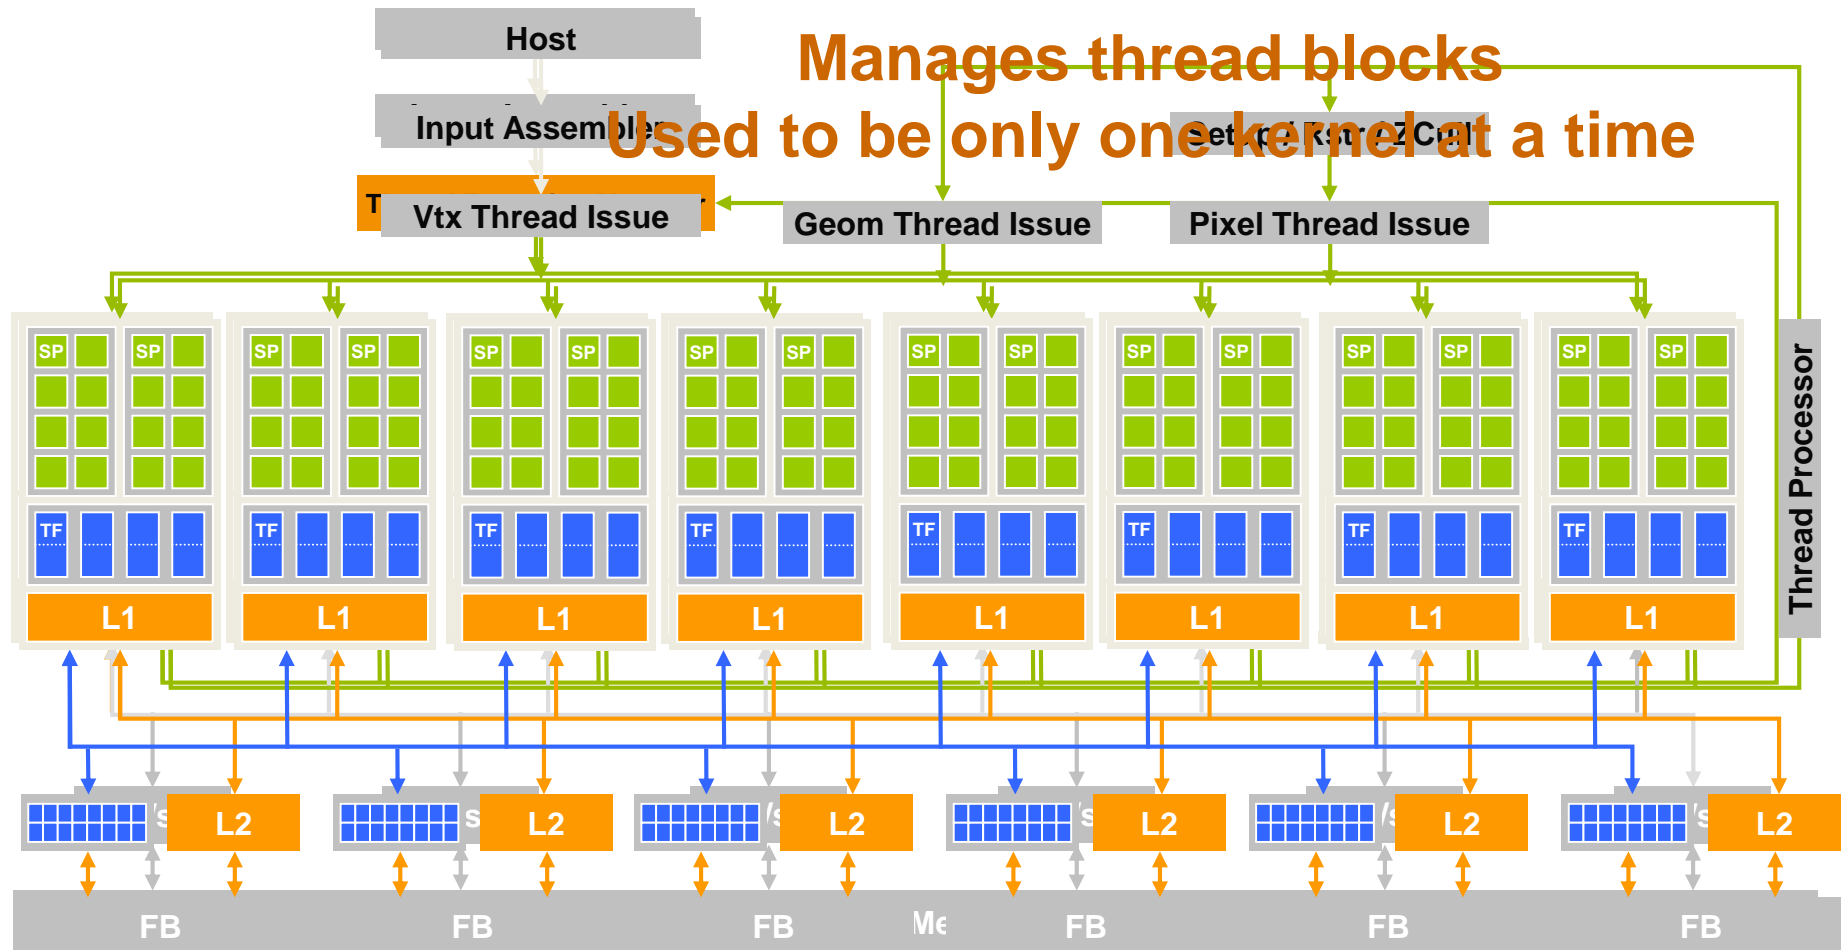
Additional information on “supplemental slides” and at
<http://graphics.stanford.edu/~kayvonf/gblog>




Make the Compute Core The Focus of the Architecture

62

- The future of GPUs is programmable processing
- After build the architecture around the processor



Old-Gen GPU Architecture: *Fermi*




- 3 billion transistors
- Over 2x the cores (512 total)
- ~2x the memory bandwidth
- **L1 and L2 caches**
- 8x the peak DP performance
- ECC
- C++
- **Announced Sept. 2009**



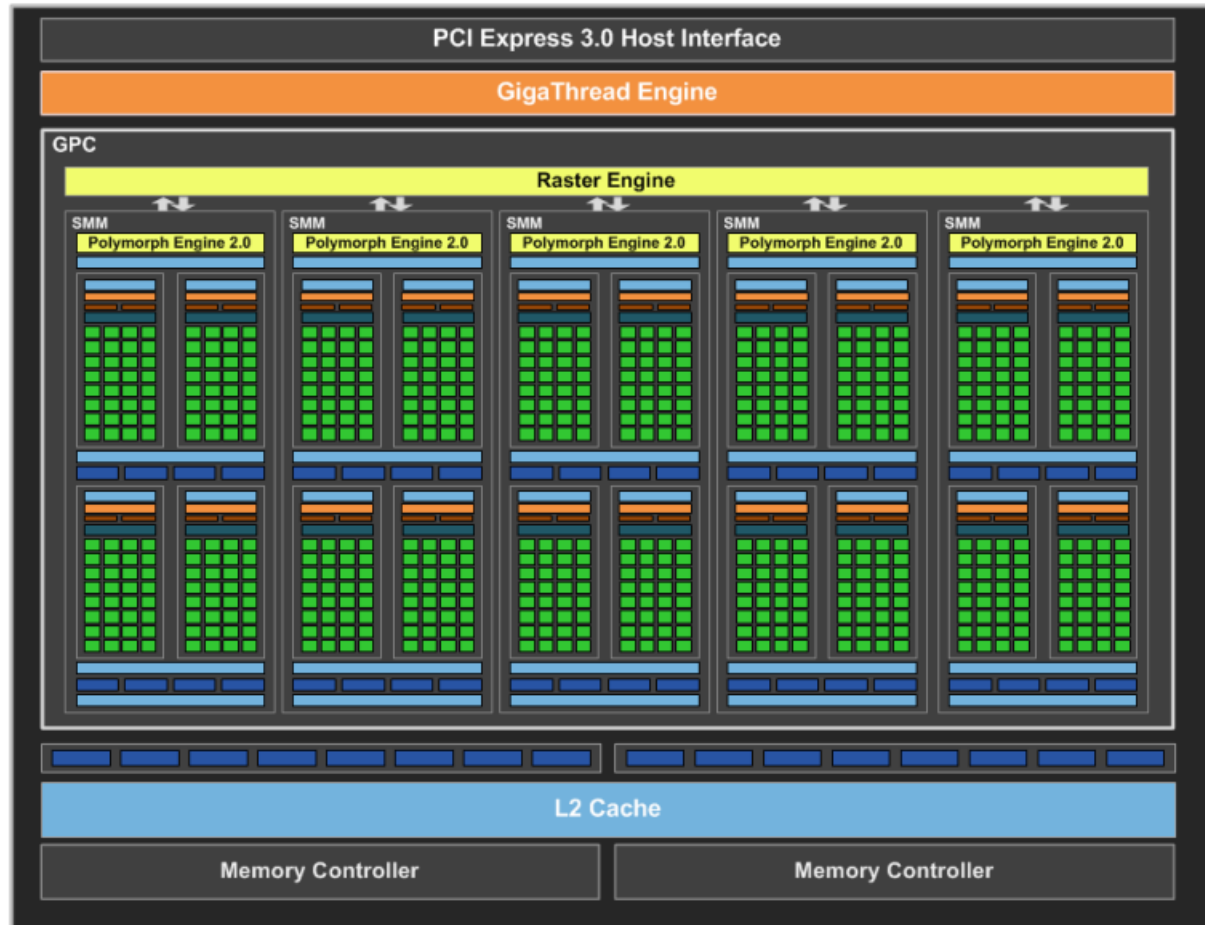
Newer GPU Architecture: *Maxwell*

64




Newer GPU Architecture: *Maxwell*

65



Even newer: Pascal




Even newer: Pascal

67



Even newer: **Pascal**


68




AMD similar but different

[Exclusive] World's First "Hawaii" Block Diagram

- Reverse-engineered Silicon Description of GCN Families
- A full Hawaii consists of 48 CUs equal to 3072 SPs / 192 TMUs



Pitcairn (212mm²)



Tahiti (352mm²)



Hawaii (438mm²)



